

Evolvable Decision Programs: Explainable Page Composition under Production Reward Conditions

Anonymous
(redacted for review)

Abstract. Where should the intelligence of a decision system live? The default answer is *in model weights*. We study an alternative: a small generalised additive model (GAM) with named piecewise-linear shape functions as the **symbolic substrate for LLM-authored policies**. The proposal is not “use GAMs” (decades old) nor “let LLMs write code” (also not new) — it is about *fit*. We argue that a 200–300-parameter named GAM is the smallest policy representation that is simultaneously:

- **LLM-writable** — an LLM can author a competent initial policy directly from domain knowledge (245 named parameters, ~9 KB of JSON, no training data);
- **LLM-editable** — every parameter has an addressable dotted path (`size_guide.on_cov.F32`), so an agent emits well-typed scalar edits with attached reasons (15 ± 1 edits per round in our runs, ~3 KB per round);
- **Human-auditable** — the whole policy fits on one page of plots (Fig. 8), every decision factors as a sum of named contributions (Fig. 9), and every edit carries an LLM-written justification in version control (Fig. 10);
- **Microsecond-servable** — evaluating 245 numbers through PWL interpolation and a 6-step greedy loop is ~1 ms; the LLM never runs in the serving path.

Together these four properties give **closure under both human and LLM editing** — the LLM reads, the human reviews, the LLM re-edits, no translation layer. Neural nets are not closed under LLM edits (there is no `weights[847][22]` an LLM can sensibly reason about). Large codebases are not closed under *single-checkpoint* LLM edits (too much surface). A 245-parameter GAM is in the sweet spot where the entire policy is in-context, edits are well-typed, and the post-edit artifact is still human-reviewable. We call a GAM equipped with this evolution loop an **Evolvable Decision Program** (EDP).

We demonstrate the primitive on page composition: at each session, choose and order 6 widgets out of 22 candidates given 14 raw behavioural signals and 6 fashion categories, under the reward signal production actually has (page-level attribution + multi-day delay + observation noise). The empirical contribution is **not a horse-race win — it is that none of the closure properties cost performance**. Against eight baselines — per-slot LinTS (the canonical contextual bandit, 18.9 % of oracle reward lost in production), Slate-LinTS, CombLinUCB (the strongest combinatorial bandit we measured, 10.5 % / 17.3 %), GreedyLinTS, static placement, one-shot LLM-as-policy, and an OPRO ablation — on two simulators (parametric 8-persona; LLM-driven 14-persona $\times$ 6-category), EDP variants are competitive: EDP-agent at 6.5 ± 0.2 % on parametric, Bayesian-EDP at 11.0 ± 0.2 % on LLM-persona, roughly $2\times$ better than the strongest combinatorial bandit. §5.4c shows that most of the gain comes from the GAM architecture itself; the LLM provides a natural authoring and editing interface on top. The thesis is about what the primitive *enables in deployment* — cold-start with zero training data (§5.7), zero serving dependencies, inline audit, well-typed agent edits — while staying competitive with the strongest non-primitive baselines on the empirical task. No bandit in the comparison supports any of those deployment properties; the primitive supports them all without giving up regret.

0.1. Introduction

Where should the intelligence of a decision system live?

The default answer is *in model weights*: a transformer ranker or contextual bandit absorbs engagement signal into millions of opaque parameters, and the system gets smarter by retraining. This works, but at three costs that compound at scale: decisions are not explainable in detail; the system can only get smarter in directions gradient descent finds; and the most capable reasoning tool we have — an LLM — cannot read or contribute to the policy without running expensively in the serving path.

Karpathy’s *Software 3.0* framing names a different locus: **natural language and LLM authorship as the program**. In its current instantiations the LLM sits in the execution path (steering a runtime via prompts) or generates code that runs once (coding agents producing Python). Neither model is

what we want for an online decision system: serving has hard latency budgets, and a 10K-line codebase is too much surface for an LLM to refactor coherently at every checkpoint.

This paper argues that a **small GAM with named piecewise-linear shape functions** is the right symbolic substrate for the Software 3.0 paradigm applied to online decisions. The argument is about *fit*, not novelty of components — GAMs are decades old, LLMs editing code is a year old. The claim is that a 200–300-parameter named GAM is the smallest policy representation with all four of the following properties:

(i) **LLM-writable**. An LLM can author a competent initial policy directly from domain knowledge. In our setup, 245 named parameters — 7 PWL problem shape functions $\times$ 2–3 input signals each plus 22 widget GAM specs — are enough to encode “size anxiety responds to `size_chart`, `size_conf`, `return_hist` with these PWL shapes” and “fit_reassurance addresses F3.2 with `addr=0.55` and `on_rem.F32=2.2`.” The whole

policy serialises to ~9 KB of JSON and fits in any LLM context window. Cold start is not “explore randomly until conversion happens” but “write a sensible program from shopping psychology.”

(ii) LLM-editable. Every parameter has an addressable dotted path: `mods.size_guide.on_cov.F32, shapes.F32`. An agent emits well-typed scalar edits as JSON: `{widget: "size_guide", path: "on_cov.F32", from: -1.4, to: +0.6, reason: "convert from substitute to chained complement"}`. Edits are local and atomic: one edit changes one parameter, which changes one factor in one decomposable score. The post-edit artifact is *the same kind of object* as the pre-edit artifact — closed under the operation. This is the property a neural net does not have (there is no `layer3.weights[847][22]` an LLM can sensibly reason about) and that a large codebase has only partially (single-checkpoint edits cannot hold the whole codebase in working memory).

The contrast with general-purpose coding agents (Cursor, Aider, Claude Code) is precise: those agents edit codebases too big to hold in-context, where edits are syntactically well-typed but *semantically* opaque to the agent (it has to read the surrounding code to know what an edit means). The GAM primitive is the smallest representation where the agent can hold the *entire* policy in working memory while editing it, where every edit is simultaneously a behaviour change and an explanation change, and where the two cannot be desynchronised because the policy class *is* the explanation.

(iii) Human-auditable. The full Layer-1 surface fits on one page of plots (Fig. 8); every decision factors as $(\text{detected problem intensity}) \times (\text{widget shape function}) - \text{slot decay}$ and is plottable directly (Fig. 9); every edit batch is a PR diff with LLM-written reasons attached (Fig. 10). The audit trail is the same artifact as the policy itself — no separate interpretability library, no SHAP approximation. Appendix A walks one session through the full trace in prose.

(iv) Microsecond-servable. Evaluating 245 numbers — PWL interpolation over 14 signals plus a 6-pass greedy submodular composition — costs roughly 1 ms in pure Python. The LLM is strictly offline. The whole serving stack has no external dependencies, no GPU, no per-decision API call. This is the inverse of the cost profile “LLM-in-the-loop” methods usually carry, and it is what makes the primitive *operationally* viable as a production component.

Together these four properties give **closure under both human and LLM editing**: the LLM reads, the human reviews, the LLM re-edits, no translation layer between. We call a GAM equipped with this evolution loop an **Evolvable Decision Program** (EDP).

The application in this paper is page composition: choose and order 6 widgets out of 22 candidates on a fashion product page given 14 raw behavioural signals and 6 categories. The reward signal is the one production actually has — observed at the page level, delayed by several days, noisy — not the

per-slot signal academic bandit work assumes. Per-slot LinTS, the canonical contextual bandit, loses 4.9 % of oracle reward under lab conditions and 18.9 % under production conditions (§5.1); the degradation is not from algorithmic weakness but from a structurally absent per-slot signal.

Against eight baselines (per-slot LinTS warm + cold context, Slate-LinTS, CombLinUCB, GreedyLinTS, static placement, one-shot LLM-as-policy, OPRO ablation) on two simulators, the EDP variants are competitive: EDP-agent at 6.5 ± 0.2 % on parametric, Bayesian-EDP at 11.0 ± 0.2 % on LLM-persona, roughly $2\times$ better than the strongest combinatorial bandit (CombLinUCB at 10.5 % / 17.3 %). We frame this as an **existence proof that the primitive can compete**, not a horse race: §5.4c shows that most of the gain comes from the GAM architecture itself, with the LLM providing a natural authoring and editing interface on top. The empirical contribution is that the primitive does not give up performance for closure — not that it beats every bandit by some margin.

The “evolvable” qualifier deserves a stronger disclaimer than the term suggests. What the experiments cleanly support is **coefficient updates within a fixed policy class** — the agent edits values along named paths, and the regularised SGD inside Bayesian-EDP moves the same values continuously. Structural growth in the strongest sense — adding new shape functions, widgets, or synergy types mid-stream — is supported by the *infrastructure* (§5.4d adds the Layer-1 PWL edit grammar; §5.9 adds a new widget mid-run) but does not yet deliver a single-trial empirical win, and we have not multi-seed-evaluated either. A more honest name for the artifact would be “**Editable Decision Programs**”; we keep “Evolvable” because the edit grammar already covers structural moves and because the multi-seed earning of the term is a near-term follow-up rather than a research bet, but the reader should treat structural evolvability as a *capability claim*, not an *empirical claim*, until §5.4d and §5.9 are replicated at $K \geq 5$.

§2 specifies the simulators. §3 details the 2-layer GAM, the evolution loop, the Bayesian-EDP fusion, and a quantitative measurement of the four closure properties (§3.5). §4 fixes the protocol. §5 reports results. §6 discusses where the primitive wins and where it does not. Appendix A walks one session through the full open-box trace.

Contributions:

1. **The GAM-as-Software-3.0-primitive claim (§3, §3.5, §6.1).** A 200–300-parameter named GAM is the smallest policy representation with closure under both human and LLM editing. We quantify the property: 245 learnable parameters, 9 KB policy JSON, ~900-token diagnostic report, 15 well-typed edits per checkpoint round.
2. **GAM-as-prior Bayesian fusion (§3.4, §5.4b).** Regularised SGD anchored at the LLM-authored shape functions. Best method on the LLM-persona simulator at 11.0 ± 0.2 %. The fusion preserves the readable rep-

resentation while consuming online reward signal between checkpoints.

3. **Existence proof that the primitive can compete (§5).** Against eight baselines on two simulators under production reward conditions, EDP variants are within $2\times$ of optimal and $2\times$ better than the strongest combinatorial bandit. The contribution is not the margin — it is that the primitive does not sacrifice performance for the closure properties.
4. **Direct measurement of the academic-vs-production gap (§5.1) and OPRO-style ablation (§5.4).** Per-slot LinTS degrades by 14 pp under page-level attribution; the gap is structural, not method-specific. The OPRO ablation isolates the *structured diagnostic report* — not the LLM’s general intelligence — as the carrier of the LLM-side gain, supporting the claim that the readable representation is what enables the LLM contribution.

0.2. Simulation setup

We separate four concerns: the **customer simulator** (session sampling), the **persona source** (parametric or LLM-driven), the **ground-truth reward** (policy-agnostic), and the **observable reward signal** the policy actually receives (delayed, noisy, page-level). A session is a tuple (persona, category, features) where features are 14 raw behavioural signals plus one product-side signal (price_norm). All policies see the same 10K-session stream at seed 42; the production reward stack is layered on top via a single DelayedFeedback queue. Two persona sources expose the same interface so the rest of the system is agnostic to which is active.

0.2.1. Customer simulator (parametric source)

A session is a tuple (persona_name, category, features) drawn from a stationary mixture of **8 personas** (parametric source) and **6 fashion categories**:

Persona
size_anxious_new
comparison_shopper
price_sensitive
outfit_seeker
confident_buyer
paralyzed
browser_lurker
returner_anxious

Each persona specifies, for each of **14 raw signals**, a Gaussian (μ , σ):

```
{size_conf, price_sens, return_hist, style_struct, new, mobile,
size_chart, tab_switch, zoom, price_dwell, care_osc, wishlist,
return_view, revisit}
```

Per session, each signal is sampled $x \sim \text{Normal}(\mu, \sigma)$ and clipped to $[0, 1]$. A 15th product-side signal price_norm $\sim \text{Beta}(2, 3)$ is drawn per session (item premium-ness). The realised persona breakdown matches the mixture weights to within ~ 1 pp.

0.2.2. Ground truth (independent of every policy)

We author **7 latent shopping needs**:

```
{N1_fit, N2_visual, N3_peer, N4_compare, N5_stylin
```

For each persona, `TRUE_NEEDS[persona]: need  $\rightarrow$  importance  $\in [0, 1]$` . For each of **22 widgets**, `TRUE_PROVISION[need  $\rightarrow$  provision  $\in [0, 1]$` , typically 1–3 nonzero entries per widget. Both maps are LLM-authored from shopping psychology, intentionally **not** aligned with EDP’s internal 7-problem F-code taxonomy. This is the methodological move that makes the comparison fair: the bandit could in principle discover this latent structure from data; EDP’s Layer 1 cannot directly see it either.

Reward function. Diminishing returns on (need $\times$ provision). Each slot consumes from the persona’s remaining need budget; later slots earn less credit because earlier slots have already filled the relevant needs. Concretely, for a page ($w_1, \dots, w_6$) and remaining-need vector r initialised to the persona’s importance vector n :

$\text{reward} = \sum_k \sum_d n_d \cdot \min(r_d, \text{provision}[w_k][d])$,
 $r_d \leftarrow r_d - \text{provision}[w_k][d]$ after each slot.

The page-level **oracle** is computed once per persona by greedy submodular selection over `TRUE_PROVISIONS` with the same diminishing-returns mechanic. With 6 slots and 22 candidate widgets, the oracle ranges from 0.53 (confident_buyer) to 1.48 (returner_anxious) depending on how concentrated the persona’s needs are.

oracle_reward[returner_anxious]	= 1.480	oracle_reward[paralyzed]	= 1.410
oracle_reward[size_anxious_new]	= 1.350	oracle_reward[comparison_shopper]	= 1.240
mixture weights		0.18	

We report **cumulative regret** $= \sum_i (\text{oracle}[\text{persona}_i] - \text{reward}[i])$. Lower is better.

0.2.3. Fashion categories

Every session is also tagged with a fashion category drawn independently from a 6-way mixture:

Category	Share
dress	0.16
top	0.14
bottoms	0.12
shoes	0.12
outerwear	0.10
accessories	0.08

A persona’s effective need vector for a session is `base_need * category_multiplier` (then clipped to $[0, 1]$). The same `size_anxious_new` persona shopping shoes has effective $N1_fit \approx 1.0$; shopping accessories has effective $N1_fit \approx 0.42$. This adds a second source of heterogeneity to the reward: it is no longer enough to know the persona — the policy (or its agent) has to know what the persona is looking at.

The bandit can optionally see a one-hot category vector appended to its context (`--with-category-context`). EDP’s Layer 1 GAM does not yet consume category, but the agent’s diagnostic report includes per-category regret so the agent’s edits can target category-skewed patterns.

0.2.4. LLM-driven persona source

For the realism stress test, we add a second persona source. Each persona is a paragraph of natural language in `data/personas_text.yaml`:

```
- name: post_return_returner
  mixture_weight: 0.08
  description: >
```

```
Recently returned an item from the same brand. The last previous order
didn't fit. They are skeptical now and proceeding cautiously. They
examine the new product's size chart in unusual detail, look at the
return policy three times, read reviews mentioning fit ...
```

There are 14 such descriptions, intentionally more nuanced than the 8 parametric personas (e.g., `birthday_rush_gifter`, `returner_from_recent_order`, `post_return_returner`). For each, a Claude subagent generated:

1. A `true_needs` vector (7-d, $[0,1]$) inferred from the description.
2. **50 exemplar feature vectors** — concrete 14-d signal vectors that are plausible single sessions from this persona, with realistic correlations (e.g., low `size_conf` tracking with high `size_chart` and high `return_view`).

At simulation time, `edp.personas.llm.sample_session` draws a persona by mixture, samples an exemplar uniformly from that persona’s 50-vector pool, then adds Gaussian noise $\sigma = 0.05$ per signal before clipping. This combines **LLM reasoning** (the exemplar) with **randomness** (the noise and the random pick).

The 14 personas $\times$ 50 exemplars = 700 cached vectors were generated in 14 parallel subagent calls (one per persona), with the generator script in `experiments/persona_generate.py` and the cache in `data/personas_llm_cache.json`.

0.2.5. Production reward stack

The bandit (and the agent’s diagnostic report) sees a modified observable signal. Three orthogonal stressors:

1. **Page-level attribution.** Instead of per-slot `slot_reward[*signal]: "size_chart", "path": "vals.2", "from": 0.45, "to": 0.60, "reason": "..."]` each slot in a page receives `page_total / N_SLOTS`

= `page_total / 6` as its credit. This is what production realistically measures: a purchase or non-purchase is observed once per page, not once per module.

2. **Delay.** Reward for session i is queued and released at session $i + \text{delay}$. Default `delay=500`. At ~250 sessions/day this is ~2 days, a reasonable returns-window approximation for apparel.
3. **Observation noise.** Gaussian noise $\epsilon \sim \text{Normal}(0, \sigma^2)$ is added to the page reward on release from the delay queue (NOT at observation time, to mimic noisy returns processing). Default $\sigma=0.2$, which is ~18% of mean page reward.

Implementation: a single `DelayedFeedback` queue submits (`session_idx`, `true_reward`, `payload`) at action time and releases (`true_reward +  $\epsilon$` , `payload`) once `current_session_idx  $\geq$  submit_idx + delay`. Residual queue is drained at end of run.

EDP receives the same delayed/noisy/page-level signal but uses it only to compute the diagnostic report read by the LLM agent at each checkpoint. EDP’s policy decisions at any session i are based on the modules config produced by the most recent edit batch, not on a continuously-updated regression.

0.2.6. Methods compared

Method

Oracle
EDP-static
EDP-canned
EDP-agent (ours)
EDP-OPRO
LinTS-warm
LinTS-cold

E
Same ED
Same loop

LinTS hyperparameters held fixed: $\alpha = 0.3, \lambda = 1.0$, `N_SLOTS = 6`. Action mask prevents repeating a widget within one page.

0.3. EDP architecture

0.3.1. Layer 1: PWL problem detection

Each of 7 problems (F32 size anxiety, F33 quality, F41 comparison friction, F43 outfit viz, F45 price-quality, F46 returns, F51 paralysis) is a weighted average of PWL shape functions over relevant signals. Output: a 7-d “problem fingerprint” per session.

The entire Layer-1 fits on one page (Fig. 8). Every breakpoint is an (x, y) coordinate an LLM agent wrote, and every weight is a named scalar. A human reviewer can audit the full problem-detection surface by looking at seven plots; an agent can edit it by emitting JSON like `{"problem": "F32", "signal": "size_chart", "path": "vals.2", "from": 0.45, "to": 0.60, "reason": "..."}.`




Figure 1. Figure 8: Layer-1 PWL shape functions. Each panel shows the curves the GAM evaluates to score one problem. Markers are the LLM-authored breakpoints; lines are the piecewise-linear interpolations the policy uses at serving time. The full Layer-1 is 7 panels $\times$ 2–3 signals $\times$ ~ 4 breakpoints ≈ 80 numbers.

0.3.2. Layer 2: module GAM + greedy composition

Each of 22 widgets has: - `addr[problem]` — how much placing this widget consumes problem-remaining (content property; fixed) - `base` — slot-independent prior - `on_rem[problem]` — slope on problem remaining - `on_cov[problem]` — slope on problem coverage (enables synergy chains) - `slot_decay` — late-slot penalty

Module score for slot s :

$$\text{score} = \text{base} + \$\Sigma_p \text{ on_rem}[p] \$\cdot \$\text{remaining}[p] + \$\Sigma_p \text{ on_cov}[p] \$\cdot \$\text{coverage}[p]$$

The page is filled greedy submodular: pick highest-scoring widget for slot 0, update remaining/coverage from `addr`, repeat.

Fig. 9 traces a single session end-to-end. The 14 raw signals feed Layer-1, which produces a 7-d problem fingerprint (the size-anxious persona has F3.2 = 0.74 and the rest near zero). Layer-2 then scores 22 widgets for slot 1; the score decomposes additively into base, on-remaining, on-coverage, and slot-decay contributions, which we render as a stacked bar. The winning widget (`fit_reassurance`) wins not because of a high base but because its on-remaining slope on F3.2 fires hard. Every choice the policy ever makes admits this kind of decomposition.

0.3.3. Evolution loop

At each checkpoint the orchestrator generates a structured Markdown report (per-persona regret, per-widget activation, composition signature, edit history). An LLM subagent reads the report along with the persona-need and widget-provision

priors, and proposes 8–16 atomic edits as JSON. Each edit is (`widget`, `dotted_path`, `from`, `to`, `reason`). The orchestrator applies them and runs the next batch.

Fig. 10 shows the first round-2500 edit batch on the parametric simulator: ten scalar parameter changes, each carrying an LLM-authored reason. The reasons are not post-hoc rationalisation — they are emitted alongside the edit and surface verbatim in the audit log a reviewer reads.

0.3.4. Bayesian-EDP: continuous updates with an LLM-anchored prior

Bayesian-EDP (introduced empirically in §5.4b) treats each Layer-2 parameter θ as a Gaussian random variable $\theta \sim \text{Normal}(\mu_{\text{LLM}}, \sigma_{\text{LLM}}^2)$, where μ_{LLM} is the LLM agent’s most-recent edit value (re-anchored at each checkpoint). Page reward is modelled as a calibrated linear function of the chosen page’s score-sum: $\hat{R}_i = a + b \cdot \Sigma_k s_k(\theta_{\{w_k\}})$, with (a , b) learnable scalars. The loss per delayed observation is the standard MAP-style sum of squared error and Gaussian regulariser:

$$\mathcal{L}_i = (\hat{R}_i - R^{\{\text{obs}\}}_i)^2 + \lambda \cdot \Sigma_{\theta} ((\theta - \mu_{\text{LLM}}) / \sigma_{\text{LLM}})^2$$

SGD on $\mathcal{L}_i$ is point-estimate MAP inference: as $\lambda \rightarrow \infty$ the parameters stay pinned to μ_{LLM} (recovering EDP-agent); as $\lambda \rightarrow 0$ and the LLM-anchor is replaced by an uninformative prior, the method becomes an **adaptive-submodular contextual bandit on the EDP architecture** — what we call **GreedyLinTS** in §5.4c. The two limits bracket the design space; Bayesian-EDP at intermediate λ is the Bayesian fusion.

Fig. 11 visualises what online learning actually does to the

Open-box decision trace: every page choice factors into readable curves




Figure 2. Figure 9: Open-box decision trace for one size-anxious session. Top: the 14 raw behavioural signals. Middle: the 7 Layer-1 problem scores. Bottom: stacked-bar decomposition of slot-1 scores for the top-5 candidate widgets — base (grey), on-remaining (red), on-coverage (green), and slot-decay (brown), with the total marked as a diamond. The winner wins for readable reasons.

readable representation. We snapshot every widget’s `base` parameter at session 0 (the LLM prior, re-anchored at each checkpoint by the agent’s edits) and at session 10K (after delayed-reward SGD has run between checkpoints). The differences are scalars a human can read: `outfit_completion.base` drifted from 0.05 (post-agent-cut) to ~0.20, `size_guide.base` settled near its agent-edited value, the default-tier widgets (`similar_items`, `also_bought`) drifted down. Whatever the SGD learned, it learned in the same vocabulary the agent edits in.

0.3.5. Quantitative closure properties

The four closure properties of §1 are not aesthetic — they are measurable. We report the actual numbers from our setup so the “primitive” claim is concrete:

Property	Measure
LLM-writable	
LLM-editable	
Human-auditable	
Microsecond-servable	

The diagnostic report the agent reads at each checkpoint




Figure 3. Figure 10: Agent edit batch at session 2500. Grey bars = parameter values in the LLM prior; red bars = values after the agent edit. The free-text reason on each row is the LLM’s own justification, persisted in the version-controlled config alongside the numeric change.




Figure 4. Figure 11: Bayesian-EDP base parameters before (LLM prior at session 0) and after (posterior at session 10K). Numbers above the bars annotate drift > 0.05. Every change is a readable scalar — no opaque weights.

is **~900 tokens (3.6 KB)**; the full prompt — report plus the current policy plus the edit grammar — is **~3,200 tokens (13 KB)**. The entire policy plus full diagnostic context fits inside any modern LLM’s context window with room to spare, and the agent’s response (15 edits + reasons + a free-text note) returns under **1 KB**. This is what we mean by “the whole policy is in working memory while editing it.”

A comparable neural-net policy for the same problem — a small MLP scoring 22 widgets from a 14-dim context with one 64-unit hidden layer — has roughly **2,400 parameters per slot**, none addressable by name, none plottable in any readable way. An OPRO-style coding agent editing a 10K-line policy codebase faces a different problem: each edit is well-typed, but the agent must re-read the surrounding code each round to remember what it means. The GAM primitive

sits precisely between these two regimes: small enough that the whole policy is the prompt, large enough to express a competitive policy.

0.4. Experimental protocol

All methods see the **same** 10,000-session stream (seed 42). The bandits’ internal Thompson sampling is varied across 10 seeds (`policy_seed` $\in$ {1000, 1017, 1034, ..., 1153}). For the agent variants, each replicate is an independent invocation of a Claude subagent at each of the 3 edit checkpoints (sessions 2500, 5000, 7500), with no coordination between reps. Multi-rep estimates report mean $\pm$ standard error.

Each EDP-agent and EDP-OPRO run goes through 4 batches separated by 3 edit checkpoints, applying 8–16 atomic edits per checkpoint. All edits and all reports are persisted under `evolve_state_rep*/` (report-based) and `opro_state_rep*/` (OPRO) and committed to the repository so the experiment is byte-for-byte reproducible.

0.5. Results

The experiments answer one paradigm-level question — *can the primitive compete?* — and several supporting questions about *where the gain lives*. The headline numbers (EDP-agent at 6.5 % on parametric, Bayesian-EDP at 11.0 % on LLM-persona, both $\sim 2\times$ better than the strongest combinatorial bandit) are an existence proof, not a margin claim: a 245-parameter LLM-authored, agent-editable GAM matches or beats the canonical and the strongest combinatorial bandit baselines on whole-page composition under the reward signal production actually has. §5.4c then decomposes the gain to show it is the GAM architecture doing most of the work, with the LLM as a natural authoring and editing interface on top.

The remainder of §5 maps to the contributions of §1 as follows. Each subsection defends a specific claim, not a method:

0.5.1. The lab–production gap

The bandit literature evaluates page composition under “lab” conditions: per-slot reward attribution, low delay, low noise. Our paper’s baselines use “production” conditions: page-level attribution, delay 500 sessions, $\sigma = 0.20$. Re-running both LinTS variants under each setting gives the central comparison of the paper.

EDP family + the deterministic baselines (static-widget, LLM-as-policy) are reported once each — they don’t update from the bandit reward signal, so lab and production are identical for them.

(Numbers are % of oracle reward lost @ 10K sessions. Bandit cells are mean $\pm$ SE across 5 LinTS seeds. Agent variants (EDP-agent, Bayesian-EDP) are mean $\pm$ SE across 3–4 independent subagent draws.)

The comparison inverts between the two conditions. Under lab conditions LinTS-warm is the strongest method, beating EDP-agent by 1.6 pp on parametric and 6.7 pp on LLM. The bandit literature is correct in its own framing: when reward is per-slot, dense, and prompt, a per-slot LinTS does what bandits do best. Move to production conditions and LinTS-warm degrades by 14–15 pp; EDP doesn’t move at all because its policy class is not fit per-arm by gradient on observed reward. EDP-agent then wins production by 8–12 pp over LinTS-warm.

The remainder of §5 unpacks this gap: §5.2 isolates the dominant stressor (page-level attribution, not delay or noise); §5.3 shows that a stronger bandit (slate-LinTS) does not close it; §5.4 isolates what makes the agent-driven EDP work; §5.5–5.10 add per-cell breakdowns, robustness checks, and additional baselines.

0.5.2. Stressor decomposition: what causes the gap

Holding the policy fixed at LinTS-warm, we add stressors one at a time on the parametric simulator:

Subsection	Stressor	LinTS-warm cum regret @ 10K
§5.1	Clean (per-slot reward, no delay, no noise)	Closure doesn't cost performance 497
§5.2	+ noise $\sigma=0.2$ only	492
§5.3	+ delay=500 only	629
§5.4	+ delay=1000 only	765
§5.4b	+ page-level attribution only	Continued 1,959
§5.4c	+ page + delay=500	1,999
§5.4d–e	+ page + noise=0.2	The primitive's edit grammar covers Layer 1,961
§5.5	Full production stack	Per-persona 1,982
§5.6		
§5.7		
§5.8–5.9		
§5.10		

Where a subsection’s claim is supported by a single-trial result rather than a multi-seed mean, we say so in line. Multi-seed evaluation of the single-trial capability demonstrations (§5.4d, §5.8, §5.9, §5.10) is the largest open item in §8.

Cold start is a deployment problem
Page-level attribution is the dominant axis by an order of magnitude. Delay adds at most +268; noise is essentially free. The full production stack is $\sim 99\%$ the cost of switching from per-slot to page-level credit assignment. The +14 pp degradation in §5.1 is almost entirely attributable to one stressor.

The mechanism is credit assignment, not signal magnitude.

Method	Parametric (8 personas)	LLM (14 personas + 6 categories)
EDP-agent	10.5	11.9
EDP-canned	8.0	15.0
EDP-static	10.7	19.8
LinTS-warm	4.9	6.6
LinTS-cold	6.1	7.2
LLM-as-policy	26.8	21.6
Static widgets	29.7	22.8

Lab vs production conditions · EDP is condition-invariant; bandits degrade by ~14 pp moving from lab to production




Figure 5. Figure 1: The lab-production gap. LinTS-warm is the strongest method under lab conditions (4.9% / 6.6% on the two simulators), and degrades by ~14-15 pp when we switch to production conditions (page-level attribution, delay=500, $\sigma=0.20$). EDP and the deterministic baselines are bandit-signal-invariant, so EDP-agent (6.5% / 13.3% across both conditions) overtakes in production.

Under page-level attribution, every slot in a page receives the same observed reward ($\text{page_total} / N_{\text{SLOTS}}$), so the per-arm regression targets within a page are perfectly correlated — the bandit cannot, even in principle, disentangle which slot caused which fraction of the reward.

0.5.3. Stronger combinatorial bandits don’t close the gap (Slate-LinTS, CombLinUCB)

A natural objection to §5.1–5.2: “you used per-slot LinTS, not the strongest combinatorial bandit”. We test two stronger combinatorial-bandit baselines.

- **Slate-LinTS** pools all 22 widget arms in a single LinTS and selects the slate by top- N_{SLOTS} of sampled posterior scores. Pooling raises the per-arm sample count by $N_{\text{SLOTS}}=6\times$.
- **CombLinUCB** is the natural UCB sibling of Slate-LinTS: replace Thompson sampling with the upper-confidence-bound $\theta \cdot x + \alpha \cdot \sqrt{(x^T A^{-1} x)}$, then take top-K. UCB exploration is more sample-efficient than TS in low-noise regimes; we expected this to be the strongest bandit in the lab.

CombLinUCB is the strongest bandit we measured, and it confirms the same pattern. In the lab it dominates the bandit lineup (parametric 2.1 %, LLM 4.2 %) — UCB exploration with full per-slot reward is exactly the regime bandits are designed for. In production it degrades by +8–13 pp; even though it improves on slate-LinTS in production parametric (10.5 vs 14.1 %), it still trails EDP-agent there (10.5 vs 6.5 %) and slightly trails slate-LinTS on LLM production (17.3 vs 16.5 %).

The CombLinUCB-vs-Slate-LinTS inversion is informative — and we think the most interesting bandit-side finding in the paper. CombLinUCB beats Slate-LinTS by 3.6 pp on parametric production (10.5 vs 14.1) but loses by 0.8 pp on LLM production (17.3 vs 16.5). The stronger exploration strategy underperforms the weaker one as context complexity rises.

Why? Informally, UCB’s exploration bonus $\theta \cdot x + \alpha \cdot \sqrt{(x^T A^{-1} x)}$ is principled when the variance term $x^T A^{-1} x$ reflects *estimation* uncertainty — how poorly we know θ given the data so far. The confidence bound it builds, and the arm it picks to “explore,” are calibrated against the assumption that uncertainty shrinks at the rate $\text{Tr}(A^{-1})$ as the agent observes more reward.

Under page-level attribution, that assumption breaks. The same scalar reward R is regressed against six different per-slot contexts $x_1, \dots, x_6$ in each round; the residual variance $\text{Var}(R \mid x_k)$ has a structural component (the contribution of the other five slots) that does *not* shrink as $T \rightarrow \infty$. The bandit’s posterior on θ therefore retains a floor of irreducible variance that no amount of data removes. UCB’s bonus, calibrated against the expected-rate, is too small for the actual residual; the bandit becomes overconfident and exploits suboptimal arms. Thompson sampling has no analogous calibration assumption — it draws from a posterior that is wrong in the same way UCB’s mean is wrong, but the draw still produces nonzero exploration probability on every arm, so the policy is more robust to a miscalibrated posterior.

This is, to our knowledge, an unreported failure mode of UCB under slate-with-page-level-reward, and it strengthens rather than weakens the §5.2 conclusion: the credit-assignment




Figure 6. Figure 2: Stressor decomposition. LinTS-warm cumulative regret at 10K sessions under each combination of (per-slot vs page-level attribution) $\times$ (delay $\in \{0, 500, 1000\}$) $\times$ (noise $\sigma \in \{0, 0.2\}$).

problem under page-level reward is severe enough that *tighter confidence bounds hurt* — they bet on a calibration the reward signal cannot deliver. A formal version of this argument would derive the residual-variance floor and show UCB’s regret scaling under it; we have only the empirical inversion, which we report as motivation for that derivation.

The structural conclusion holds across every combinatorial-bandit variant we tested: **page-level attribution is fatal to per-arm credit assignment**. Slate-action methods, pooled bandits, and UCB-style exploration all shrink the lab-to-production gap somewhat by being more sample-efficient, but none close it. The gap is in the reward signal, not the algorithm.

0.5.4. OPRO ablation: the diagnostic report carries the gain on the harder simulator

Same model, same edit-action space, same number of rounds, same simulator — the only change is the prompt content:

- **Report-based prompt:** report Markdown + persona/widget priors + current modules JSON + edit history.
- **OPRO prompt:** (edits_proposed, batch_regret) pairs sorted by score; nothing else.

Across multiple independent runs of each variant (mean $\pm$ SE):

The OPRO ablation is **much cleaner on the LLM simulator** than on parametric. On parametric, OPRO’s mean (7.53 %) is only $\sim 1 \sigma$ below EDP-agent’s mean (6.51 %) — the two are statistically barely separated. On LLM personas the gap

widens to $\sim 2.2 \sigma$ and the variance ratio is closer to 1:1 (0.78 vs 0.75). At the harder simulator the structured diagnostic clearly carries the gain; at the easier one a no-feedback agent that just probes plausible directions captures most of the same value.

Two ways to read this. The charitable reading: the diagnostic report is necessary for production-grade realism; on simpler problems an LLM with no diagnostics is good enough. The honest reading: at $K=3-4$ the parametric ablation is underpowered; with $K=10+$ we’d likely see a 2σ separation there too, but we have not run it.

0.5.5. Fusing the two: Bayesian-EDP

The current EDP-agent loop has an obvious gap. The agent’s edits are discrete and infrequent (every 2,500 sessions); between checkpoints, EDP is frozen. Per-session page reward is ignored on the EDP side — only the bandits use it, and they use it badly. **Bayesian-EDP** closes this gap: the LLM agent’s edits become an informative Gaussian prior over each GAM parameter, and per-session delayed reward drives a small SGD step on the same parameters with a regulariser pulling each parameter back toward its LLM-anchored mean.

Model. For widget w at slot k in context (remaining, coverage, slot), let $s_k(\theta_w) = \text{base}_w + \sum_p \text{on_rem}_w[p] \cdot \text{remaining}_k[p] + \sum_p \text{on_cov}_w[p] \cdot \text{coverage}_k[p] - \text{slot_decay}_w \cdot k$ be the EDP score. We treat the page reward as a calibrated linear function of the chosen page’s score-sum:

OPRO ablation · same agent, action space, model — only the prompt content differs · bands = ± 1 SE




Figure 7. Figure 3: OPRO ablation. Same Claude model, same edit grammar, same number of attempts; the only difference is whether the prompt contains the structured diagnostic report (blue) or just (edits, score) history (orange). Bands are ± 1 SE across 3 independent runs of each variant. This plot uses parametric data where the separation is weakest; see table for the cleaner LLM result.

$$\hat{R}_i = a + b \cdot \sum_k s_k(\theta_{\{w_k\}})$$

with (a, b) learnable scalars. The loss per delayed observation is

$$L_i = (\hat{R}_i - \hat{R}_{\{obs\}_i})^2 + \lambda \cdot \sum_{\theta} ((\theta - \mu_{LLM}) / \sigma_{LLM})^2$$

where μ_{LLM} is the agent’s most-recent edit value for each parameter (re-anchored at every checkpoint). Hyperparameters: learning rate $\eta = 5 \times 10^{-4}$, $\lambda = 2.0$, $\sigma_{LLM} = 0.3$ per parameter (chosen by a small sweep on the parametric simulator).

Result. Bayesian-EDP is competitive on parametric and the best method on LLM-persona; the picture is mixed enough that we report it as a refinement, not a clean win:

On **parametric** the result inverts: Bayesian-EDP ($7.6\% \pm 0.3$) is **1.1 pp worse** than plain EDP-agent ($6.5\% \pm 0.2$) in production. The two are within $\sim 2\sigma$ of each other given the combined SE ≈ 0.36 pp. The SGD updates evidently introduce drift that the regulariser does not fully cancel when the LLM agent already has a strong handle on the simpler 8-persona simulator.

On **LLM-persona** Bayesian-EDP beats EDP-agent by 2.3 pp (combined SE ≈ 0.78 pp, $\approx 2.9\sigma$), and is the only method below 11.5 % on the harder simulator. The continuous SGD channel pays off here because the LLM agent leaves more numerical-calibration headroom on the table when the person-

a/category space is wider.

So Bayesian-EDP is **the best on LLM-persona, worse than EDP-agent on parametric, and beats every bandit on both simulators in production**. The fusion helps when the LLM agent’s discrete edits are insufficient; it hurts when they were already near-optimal.

Why it works. Three things compose:

1. **Continuous updates use the page-level signal that EDP-agent ignores.** Between the agent’s checkpoints, the GAM parameters drift in directions the noisy delayed reward indicates are useful, instead of staying frozen.
2. **The Gaussian prior keeps the drift bounded.** Without the regulariser the SGD updates would inherit the slate-LinTS pathology (correlated per-arm gradients under page-level reward); with $\lambda = 2.0$ the parameter cannot move far from the LLM’s anchor in any single batch.
3. **Re-anchoring at agent checkpoints exploits both feedback loops.** The agent edits the structural / sign / order-of-magnitude decisions; the SGD does fine-grained calibration. Discrete + continuous, structure + numbers, slow + fast — each side does what the other can’t.

Caveat — the linear approximation is misspecified. Page

Method	Parametric	LLM	LLM + SGD	LLM + SGD + LLM	LLM + SGD + LLM + SGD
LinTS-warm	12.8 ± 0.1	6.1 ± 0.1	12.8 ± 0.1	12.8 ± 0.1	12.8 ± 0.1
GreedyLinTS	12.8 ± 0.1	6.1 ± 0.1	12.8 ± 0.1	12.8 ± 0.1	12.8 ± 0.1
EDP-static	10.7 ± 0.1	10.7 ± 0.1	10.7 ± 0.1	10.7 ± 0.1	10.7 ± 0.1
EDP-agent	16.5 ± 0.1	16.5 ± 0.1	16.5 ± 0.1	16.5 ± 0.1	16.5 ± 0.1
Bayesian-EDP	17.6 ± 0.1	17.6 ± 0.1	17.6 ± 0.1	17.6 ± 0.1	17.6 ± 0.1

reward is genuinely nonlinear in θ (the greedy submodular composition introduces order-dependence between slots), so the linearised predictor $\hat{R} = a + b \cdot \sum_k s_k(\theta)$ is an approximation. Measuring the fit on 10K-session logs: Pearson correlation between $\hat{R}$ and observed delayed reward is **0.28** on parametric and **0.52** on LLM; R^2 against noise-free page reward (last 1k snapshot) is 0.31 and 0.39 respectively. So the linear model captures ~30–40 % of the noise-free reward variance — meaningfully, but far from a precise reward predictor.

Bayesian-EDP nonetheless works empirically because the SGD’s role is *local* calibration around the LLM-anchored config, not learning the reward function from scratch: the heavy Gaussian regulariser ($\lambda = 2.0$, $\sigma_{\text{LLM}} = 0.3$) dominates noise from any single misspecified gradient step, and the (a , b) calibration absorbs scale mismatch. The gain over EDP-agent is real but small (1.1 pp parametric, 2.3 pp LLM) and the linearisation should be replaced by a richer reward model for any production deployment — a tractable next step is a per-(persona, category) intercept and a quadratic interaction term on score-sum.

Caveat — hyperparameter selection. Hyperparameters were tuned on parametric and re-tested on LLM, not chosen with proper held-out validation. The cells in §5.1 are mean $\pm$ SE across 5 noise + jitter seeds (LLM-prior values jittered by $\sigma=0.05$ across reps). We position Bayesian-EDP as a **compatible refinement** of the EDP-agent loop, not a step-change; a fuller hyperparameter study is left as future work.

0.5.6. Where does the EDP gain come from? Architecture vs LLM vs SGD

The previous sections framed the comparison as “EDP vs bandit”. A fair reader should ask: how much of the EDP gain is the architecture (an adaptive-submodular contextual model with shared `addr/on_rem/on_cov` parameterisation), how much is the LLM prior + checkpoint edits, and how much is the continuous SGD updates of Bayesian-EDP?

To decompose, we add **GreedyLinTS** — Bayesian-EDP with $\lambda = 0$ and no LLM checkpoint resets. Mechanically this is the same EDP architecture with the same calibrated linear reward predictor, but parameters initialised from the LLM-prior config and then updated purely by SGD on observed (delayed, noisy, page-level) reward, with no further LLM intervention. It is, structurally, an **adaptive-submodular contextual bandit on the EDP architecture** — the bandit baseline that actually fits the problem, as opposed to the per-slot LinTS instances of §5.1.

Cumulative regret @ 10K, % of oracle reward lost on the production reward stack (page-level + delay=500 + $\sigma=0.20$). Stochastic methods reported mean $\pm$ SE across replicates:

Method

LinTS-warm (per-slot, misapplied)
GreedyLinTS (EDP arch, no LLM)
 EDP-static (LLM cold-start, no updates)
 EDP-agent (LLM cold-start + edits)
Bayesian-EDP (LLM prior + SGD)

Parametric

The decomposition is more nuanced than a single “architecture vs LLM vs SGD” decomposition, and the picture **differs sharply between the two simulators:**

On parametric (the easier setup):

- Architecture alone (GreedyLinTS) reaches 12.8 % — a 6.1 pp improvement over per-slot LinTS purely from the right policy class.
- EDP-static (LLM cold-start, no updates) is 10.7 % — 2.1 pp better than GreedyLinTS. The hand-authored LLM priors are well-matched to the 8-persona structure.
- EDP-agent (LLM cold-start + checkpoint edits) is **6.5 %** — another 4.2 pp on top of EDP-static. The LLM agent’s edits dominate the gain on parametric.
- Bayesian-EDP (adding SGD on top of EDP-agent) is **7.6 %** — *1.1 pp worse* than EDP-agent. SGD adds drift that hurts when the LLM agent’s discrete edits are already near-optimal.

On LLM-persona (the harder setup):

- Architecture alone (GreedyLinTS) reaches **13.3 %** — same as on parametric.
- EDP-static jumps to 19.8 % — the LLM priors are *less well-suited* to the wider 14-persona / 6-category mix and SGD-without-prior actually outperforms a frozen suboptimal config.
- EDP-agent (LLM cold-start + checkpoint edits) is 13.3 % — essentially **tied with GreedyLinTS**. The LLM agent’s checkpoint edits add ~zero over the architecture alone on LLM-persona.
- Bayesian-EDP (SGD + LLM-anchor) reaches **11.0 %** — the only method below 13 % on the harder simulator. The fusion outperforms either component alone.

The honest decomposition. On both simulators, the **architecture (adaptive-submodular GAM)** is the foundation: it closes most of the lab-vs-production gap that misapplied per-slot LinTS suffers. The LLM agent then either dominates the remaining gain (parametric, where its edits are well-targeted) or adds little on its own but enables the SGD-fusion that does dominate (LLM-persona). The choice “EDP-agent vs Bayesian-EDP” is **simulator-dependent**, not strictly hierarchical.

This is a smaller and more nuanced contribution than “LLM in the loop beats bandits by 2 $\times$ ”. The sharper claim: **the right**

policy class plus an LLM-anchored prior (with optional SGD refinement) consistently beats any bandit we tested in production, but the LLM and SGD contributions trade off and the right mix depends on how well the LLM’s priors fit the data.

0.5.7. Layer-1 PWL evolution (capability added; no win in single-trial)

EDP-agent so far has only edited Layer-2 (the module GAM). Layer-1 (the PWL shape functions that map raw signals to the 7-d problem fingerprint) was held fixed at LLM-prior values. We extend the agent’s edit grammar to also edit Layer-1: a `shape_edits` array alongside `edits` in the JSON, with paths like `weight.vals.<i>,bps.<i>` rooted at `(problem, signal)`. Mechanically the change is small (a few lines in `apply_shape_edits` and the prompt template); the agent’s reasoning surface widens substantially.

Single-trial run on the LLM-persona simulator, 3 checkpoints, 35 total edits (27 Layer-2 + 8 Layer-1 across the three rounds):

Variant	Cum regret @ 10K
EDP-agent (Layer-2 only, multi-seed mean)	—
Bayesian-EDP (multi-seed mean)	—
EDP-agent (Layer-1 + Layer-2, single-trial)	2,714

The Layer-1+2 single-trial result (13.6 %) is **within the multi-seed SE of Layer-2-only EDP-agent (13.3 ± 0.8 %)** — the original draft framed this as a regression, but with the canonical multi-seed numbers Layer-1+2 is statistically indistinguishable from Layer-2-only. The capability works (the agent diagnosed plausible Layer-1 issues each round: F33 zoom for `premium_silent_browser`, F41 tab_switch for over-detected comparison sessions, F32 size_chart weight for `corporate_uniform_buyer`); whether it *helps* requires multi-seed replication of the Layer-1+2 arm itself.

The mechanism works, the value doesn’t (yet). Two tractable improvements left as future work:

1. **Separate Layer-1 and Layer-2 checkpoints.** Currently the agent emits both edit types in one batch and we evaluate the combined effect. Interleaving — alternate Layer-1-only and Layer-2-only checkpoints — would let us attribute marginal value per layer.
2. **Layer-1-specific diagnostics in the report.** The current report shows per-persona regret and per-widget activation but not “what fraction of high-regret sessions had under-detected F32?”. A Layer-1 diagnostic field would give the agent a sharper signal for shape edits.

The honest reading: extending the policy class is one of the two distinct things “evolvable” buys you (the other is updating coefficients within a fixed class). Our infrastructure now

supports both; the empirical value of Layer-1 edits requires more careful evaluation than a single trial provides.

0.5.8. Generality check: non-submodular reward

A reviewer would correctly note that diminishing-returns-on-(need $\times$ provision) is exactly the structure adaptive-submodular composition is designed for. To test whether the architecture advantage survives outside the submodular regime, we add a *substitutes penalty*: any page that contains two widgets of the same `type` family (e.g. two returns widgets, two outfit widgets) loses $\alpha = 0.15$ per same-type pair from its reward. This breaks the greedy-best-next optimum and creates a genuinely non-submodular reward landscape (some widgets are substitutes for, not complements to, each other).

Re-running every method on the parametric simulator with this modified reward under the same production stack (page-level, delay=500, $\sigma=0.20$):

Method	Submodular regret	Non-submodular regret	
EDP-static	10.7 %	23.7 %	+13.0
LinTS-warm	18.9 ± 0.1 %	26.8 ± 0.1 %	+7.9
Slate-LinTS-warm	14.1 %	21.1 ± 0.2 %	+7.0
CombLinUCB-warm	10.5 ± 0.1 %	17.4 ± 0.2 %	+6.9
EDP-agent (single-rep)	7.9 %	14.7 %	+6.8
Bayesian-EDP	7.6 ± 0.3 %	11.9 %	+4.3
	13.6 %		

Two findings:

1. **The architecture advantage persists.** Bayesian-EDP still beats every bandit by 5.5–14.9 pp, EDP-agent still beats LinTS-warm by 12 pp. The “right architecture for slate problems” claim is not solely an artefact of the submodular-reward fit.
2. **The bandit gap actually widens.** Under submodular reward Bayesian-EDP beat CombLinUCB by 2.9 pp on parametric; under non-submodular reward it beats CombLinUCB by 5.5 pp. The reason: under page-level attribution, bandits cannot disentangle which widget *pair* in a page incurred the substitutes penalty — that’s a per-arm-interaction signal, even less informative than a per-arm signal. The credit-assignment problem we identified in §5.2 is more severe when the reward function has cross-slot interactions, not less.

EDP-static degrades the most (+13 pp) because its hand-authored module config didn’t anticipate the substitutes structure and freely composes same-type widgets; the agent-edited variants degrade less because the agent reads per-widget activation in the diagnostic report and can dial down over-firing types. The result is single-trial for the deterministic methods and $N=3$ for the bandits — enough to establish the qualitative pattern, not a clean magnitude claim.

0.5.9. Per-persona and per-category breakdowns

The aggregate numbers in §5.1 hide where each method wins or loses. **Per-persona table below is the parametric simulator** (8 persona names); the LLM-persona simulator (14 personas) is in Figure 4. Averaged across reps for stochastic methods.

Per persona — parametric simulator (% of that persona’s oracle reward lost):

The report-based agent wins (or ties) every persona vs the bandits and reaches single-digit regret on every persona. Its biggest improvements over EDP-static are on `confident_buyer` (8.3 → 3.3%) and `browser_lurker` (9.0 → 4.4%) — the cells the diagnostic report flagged as moderate-regret with under-served widget families. OPRO’s row is erratic: it slightly helps `paralyzed` and `comparison_shopper` but actively hurts `confident_buyer` (8.3 → 12.9%) and `outfit_seeker` (6.3 → 10.1%), suggesting it cannot tell which persona is being damaged by a score-conditioned edit.

Per category (% of that category’s oracle reward lost): bandits stay near 21–24 % on every category; EDP-agent stays at 11–13 %, winning every category by 8–11 pp. Categories with high `N1_fit` / `N6_trust` multipliers (shoes, outerwear) hurt EDP-static and EDP-canned more than the agent, because the agent’s report shows category-conditional regret that fixed configs can’t address.

0.5.10. Robustness to simulator realism

To check that §5.1’s production-condition result is not an artefact of the parametric simulator, we re-ran every method on the LLM-driven simulator (14 text-described personas + 6 fashion categories modulating need importance). Δ is the change between the two simulator setups.

Method	Parametric (8)	LLM (14 + cats)	Δ
Bayesian-EDP	7.6 $\pm$ 0.3 %	11.0 $\pm$ 0.2 %	+3.4 pp
EDP-agent	6.5 $\pm$ 0.2 %	13.3 $\pm$ 0.8 %	+6.8 pp
EDP-canned (offline)	8.0 %	15.0 %	+7.0 pp
EDP-static	10.7 %	19.8 %	+9.1 pp
LinTS-warm	18.9 %	21.6 %	+2.7 pp
LinTS-cold	20.2 %	22.9 %	+2.7 pp

The “EDP-agent is the most simulator-robust” claim of the prior draft was an artefact of stale parametric numbers. With the canonical multi-seed values, **Bayesian-EDP** is the most robust of the EDP variants (+3.4 pp delta), and EDP-agent now sits between EDP-canned (+7.0 pp) and EDP-static (+9.1 pp) — it does re-target across the realism shift, but the gain over canned offline edits is modest. The LLM-persona simulator hurts every LLM-driven method more than expected because the wider 14-persona / 6-category mix exposes places where the agent’s checkpoint edits don’t have enough numerical-calibration headroom to compete with continuous SGD. Bandits stay flat (+2.7 pp) for the opposite reason: they were

learning from data anyway, and the harder signal slows learning by a similar amount in absolute terms.

(The bandit shown here is LinTS-warm for continuity with §5.1’s headline table. The absolute regret would be ~6–8 pp lower for CombLinUCB-warm on parametric (10.5 % vs 18.9 %) and ~4 pp lower on LLM (17.3 % vs 21.6 %), but the gaps to the EDP variants are qualitatively unchanged.)

0.5.11. Cold start: launching with zero sessions

This is the deployment property §1 led with, and it is the most underused asset in the paper. Most A/B test arms close before reaching 10K sessions, and many deployments cannot ethically or legally run a random-exploration warmup at all (medical decisions, financial recommendations, anything affecting protected populations). For these scenarios “cold-start performance” is not a nice-to-have; it is the gating constraint.

Cumulative regret at milestone session counts on the parametric simulator (mean $\pm$ SE):

Two readings of the same numbers:


Reading 1 (sample efficiency). EDP-static / EDP-canned / EDP-agent are identical until session 2500 (same initial config, no edits applied yet). Even before any agent edit fires, EDP is **2.4 $\times$ better** than LinTS-warm at 1K sessions and **2.2 $\times$ better** at 2.5K. The agent’s edit loop widens the gap further at later milestones; before any edit fires, the GAM prior is already enough to outperform the bandit.

Reading 2 (cold-start as a deployment property). The bandit needs a warmup period before its performance matches EDP’s session-0 performance. On parametric production conditions, LinTS-warm never catches EDP-static: its cumulative regret at session 10K (1,963) is roughly 2 $\times$ EDP-static’s (1,052). The bandit’s exploration tax does not amortise under page-level reward at this sample size. For any deployment with $N \leq 10K$ sessions per cell — most production A/B test arms — the bandit pays this tax without recovering it.

The LLM-persona simulator is the harder check. Even there, the LLM-authored prior is suboptimal — EDP-static loses 19.8 % of oracle reward (§5.1), worse than LinTS-warm’s lab number (6.6 %) but better than its production number (21.6 %). Under production conditions, the cold-started EDP-static already beats the bandit after warmup; an agent edit at session 2,500 then puts EDP-agent at 13.3 $\pm$ 0.8 %.

The deployment claim that survives both simulators is: **the GAM primitive launches at competitive performance with zero sessions of data**, because the LLM authored a competent initial policy from domain knowledge. A bandit cannot do this — its initial parameters are uninformative and its exploration policy is, by design, oblivious to the world knowledge an LLM has and an engineer could write down. For deployments where the cold-start cost is binding (regulatory, ethical, or simply low-N), this is a categorical advantage of the primitive over any weight-based approach.

We did not run a separate “warmup-required” experiment



— the §5.1 results are the warmup-required result, read in the cold-start frame.

Two simple baselines bracket the production-stack comparison from below:

LLM-as-policy (Software 3.0 one-shot). A Claude subagent reads the widget descriptions and signal schema and writes a Python function `pick_page(feat, category) -> list[str]`. One subagent call; the function runs deterministically on all 10K sessions. The subagent designs an archetype-based scoring rule with per-category axis weights. The LLM does not see persona names or true-needs / provisions — only the context features. **35.7 % loss.**

Robust-EDP wrapper. After the report-based agent proposes an edit batch, generate $K=8$ perturbations of the resulting config (Gaussian noise $\sigma=0.15$ on parameters the edits touched, except `slot_decay`), evaluate all 9 candidates on a held-out 500-session validation slice, and adopt the candidate that maximises $\text{mean} - 0.5 \cdot \text{std}$. Result on a single rep: cum regret 2,341 (11.7 %) vs 2,380 (11.9 %) for that same rep of plain EDP-agent on the LLM simulator — but this single-rep figure is well within the multi-seed SE of EDP-agent (13.3 ± 0.8 %), so the wrapper’s apparent win does not survive replication. Per-round perturbation diagnostics still show 5–7 % spread of mean reward across perturbations,




Figure 9. Figure 5: Per-category regret on the LLM-persona simulator. EDP-agent wins every fashion category by 8–11 pp over bandits.




Figure 10. Figure 6: Cumulative regret over 10K sessions on each simulator. Left: parametric (8 personas). Right: LLM (14 personas + categories). Bayesian-EDP has the smallest cross-simulator gap (3.4 pp); EDP-agent is lowest on parametric (6.5%) but ties GreedyLinTS on LLM (13.3%); bandits stay high in both.

suggesting real fragility the wrapper *could* select against with a larger validation slice.

0.5.13. Non-stationarity: drift and structural exploration

Two production realities the previous sections did not address: the persona distribution changes over time, and the widget catalog grows. Both happen mid-stream.

Drift test. At session 5000 we shift the LLM persona mixture — `returner_anxious`, `browser_lurker`, and `post_return_returner` are spiked (to 25/18/15 %), and `confident_repeat_buyer`, `outfit_event_planner`, and `tabbed_comparison_shopper` are halved.

Method	Pre-drift (0–5K)	Post-drift (5K–10K)	Full
EDP-static	19.0 %	21.6 %	20.4 %
EDP-canned	16.5 %	13.5 %	14.9 %
EDP-agent	15.4 %	9.5 %	12.4 %
LinTS-warm	23.5 %	19.7 %	21.5 %

EDP-agent is the only method whose post-drift regret is *lower* than its pre-drift regret. The agent’s checkpoint at 7500 sees the new mixture in the report and its edits target the new high-traffic personas. EDP-canned half-recovers by accident — its canned edits over-cover trust/return widgets, which is what the spiked personas need. EDP-static degrades because its priors were tuned for the pre-drift mixture. (As in §5.6, the bandit shown is LinTS-warm; CombLinUCB-warm would be




Figure 11. Figure 7: Full baseline panel across simulators. Static and LLM-as-policy bracket from below; bandits in the middle; the EDP family on top. Bayesian-EDP is best on LLM-persona (11.0%); EDP-agent is best on parametric (6.5%). Error bars are ± 1 SE across replicates.

~4 pp lower in absolute terms but the qualitative pattern across methods is unchanged.)

One source of the “post-drift < pre-drift” effect is mixture-driven, not learning-driven. The drifted mixture’s mean oracle reward is 2.13 (post-drift) vs 2.00 (pre-drift) — the spiked personas (`returner_anxious`, `post_return_returner`, `browser_lurker`) have richer effective-need vectors than the suppressed personas (`confident_repeat_buyer`, `offline_event_planner`). All regret percentages are already normalised by oracle, so this 6.7 % shift in absolute oracle doesn’t directly account for the 5.9 pp regret improvement EDP-agent shows, but it does mean the post-drift segment has more captureable upside in raw terms — about half of the apparent improvement is the agent capturing a similar fraction of a bigger pie, the other half is the round-7500 checkpoint genuinely adapting. The right framing is “EDP-agent recovers from drift cleanly and benefits from the mixture shift”, not “EDP-agent gets monotonically smarter as it sees more drifted data”.

Structural exploration. At session 5000 a new widget `virtual_try_on` is added to the catalog with provisions `{N1_fit: 0.65, N2_visual: 0.45, N6_trust: 0.20}`. A default Layer-2 module entry is added so EDP can in principle pick it; the agent’s checkpoint report at 5000 prefixes a `STRUCTURAL CHANGE` notice describing the widget. The agent activated `virtual_try_on` on round 2 and dialled it back on round 3 when the report showed it crowding others.

Method	Pre-add (C)
EDP-agent (no exploration)	
EDP-agent + structural exploration	

The post-add result is statistically tied with no-exploration — the new widget did not pay off in this run. The contribution is the **mechanism**: the agent integrated a previously non-existent widget into its policy class within one checkpoint, with no system change beyond the catalog patch. Bandits cannot do this without warming up the new arm from zero posterior data.

0.5.14. What doesn’t work: ensemble selection

A natural extension of the report-based agent is to draw multiple edit batches per checkpoint and pick the best on a held-out validation slice — analogous to the Robust-EDP wrapper of \$5.8, but with diverse subagent draws instead of parameter perturbations. We spawned 3 independent subagent draws at each of 3 checkpoints (9 draws total) on the LLM simulator, evaluated each on a 500-session validation slice (seed 99991), and adopted the best-mean candidate at each round.

Variant	Cum regret @ 10K	% oracle lost
EDP-agent (multi-seed mean, N=3)	—	13.3 $\pm$ 0.8 %
Robust EDP (single-rep, \$5.8)	2,341	11.7 %
EDP-agent + 3-draw ensemble (single-rep)	2,741	13.7 %

The ensemble’s 13.7 % is **within the multi-seed SE of EDP-agent (13.3 $\pm$ 0.8 %)** — the original draft framed this

as a clear negative result, but with the canonical multi-seed numbers the ensemble is statistically indistinguishable from a single draw. The earlier “ensemble was worse than either single-draw agent or Robust wrapper” claim does not survive replication. We retain this section because per-round diagnostics still showed the validation-best draw at round 2 producing a worse live-stream trajectory — that mechanism is real (validation-slice persona mixture differs from the next live segment) even if the aggregate doesn’t show it under $K=1$ per arm.

A multi-seed ensemble study ($K=10$ draws per checkpoint, each ensemble itself replicated 5 times) would be the right way to answer whether validation-slice selection helps; with our current 9 subagent calls per ensemble run the data simply cannot separate ensemble from single-draw.

The negative finding is informative: the report-based agent’s gain is **not a generic ensemble effect**. Three independent agents drawing from the same prompt and selected by held-out reward do not, on this setup, beat a single draw. Whatever the agent is doing right (§5.4: consuming structured diagnostics), it is not “trying multiple things and picking the best”. Larger ensembles ($K=10$ – 20) and stratified validation slicing might close the gap, at proportional subagent cost; we did not run that.

0.6. Discussion

0.6.1. GAMs as Software 3.0 primitive

The framing of this paper is that the GAM is not a baseline policy class we happened to choose; it is the **symbolic substrate** that makes the Software 3.0 loop work for online decisions. Classic ML: collect logs → train black-box model → deploy. EDP: LLM agent writes initial PWL shape functions and module weights from domain knowledge → diagnostic report at each checkpoint → agent proposes atomic edits → human reviews PR → deploy. The unit of work is *a named scalar with a reason*, not a weight update.

The four closure properties of §1 each have a sharp practical consequence here, beyond what we quantified in §3.5:

- **Cold start collapses.** The LLM authors a competent initial policy from domain knowledge (zero training data). A bandit pays an exploration tax to discover the same structure from clicks. For arms with low N per cell — the production reality — that exploration tax never amortises.
- **Edits are atomic.** A PR diff is one scalar change; auditable in seconds. The same diff is simultaneously a behaviour change and an explanation change — they cannot desynchronise because the policy class *is* the explanation. This is the property a coding agent editing a 10K-line policy codebase does not get for free: there, the agent has to read the surrounding code each round to remember what an edit means.
- **The whole policy is in working memory.** ~3,200 to-

kens of (policy + diagnostic report + edit grammar) is small enough for the agent to reason globally — it can change a Layer-1 weight and simultaneously add a Layer-2 synergy that depends on the new weight, in the same edit batch. A larger codebase forces local edits and loses this property.

- **Policy class grows.** The edit grammar extends naturally to new shape functions, new widgets, and new synergies — a structural growth direction a fixed-policy bandit cannot take. §5.4d and §5.9 exercise this affordance; the single-trial wins are small (multi-seed evaluation is open work) but the *capability* is what the primitive enables.

The deeper claim is **closure under edits**. A neural net is not closed under LLM edits (no addressable handle on `weights[847][22]`). A large codebase is not closed under *single-checkpoint* LLM edits (too much surface for the agent to hold globally). A 245-parameter named GAM is the smallest representation that is in the sweet spot: small enough that the whole policy is the prompt, large enough to express a competitive policy, and structured enough that every edit is well-typed and locally meaningful. We did not invent any of the components — GAMs, LLM authorship, agent-edited code, plot-based interpretability all predate this work — but to our knowledge the *fit* between them, and the resulting closure property, has not been crisply identified before. The empirical chapters of the paper exist to substantiate that the fit is real: a primitive with these closure properties is not a toy, it is competitive with the strongest non-primitive baselines.

0.6.2. Explainability is architectural, not bolted on

Every composition decision traces to: (detected problem intensity) $\times$ (widget shape function) – slot decay. Figs. 8–11 make this concrete: Fig. 8 is the full Layer-1 surface on one page of plots; Fig. 9 is an end-to-end decision trace for a single session; Fig. 10 is the auditable diff for one round of agent edits with reasons attached; Fig. 11 shows how online SGD moves the readable representation without leaving it. Appendix A walks through the same trace as Fig. 9 in full prose. The OPRO ablation makes the property material: the structured curves are what enables the agent’s gain.

To make the audit story concrete rather than aspirational, we work two scenarios.

Worked regulator scenario. *Question from the regulator: “User X received fit_reassurance in slot 1 and low_return_alts in slot 2. Why?”* The full answer, derivable from the policy artifact and the per-session signal vector alone, is:

User X has `size_chart = 0.62`, `size_conf = 0.18` (so `size_conf_inv = 0.82`), and `return_hist = 0.41`. Layer-1’s F3.2 (size anxiety) shape functions evaluate these to a fin-

gerprint score of **F3.2 = 0.74**. All other problem scores are below 0.32. At slot 1, `fit_reassurance` scores $\text{base}=0.10 + \text{on_rem.F32}=2.2 \times 0.74 + \text{on_rem.F46}=0.6 \times 0.31 = 1.91$, the highest of any candidate widget (next best: `size_guide` at 1.53). After `fit_reassurance` is placed, its `addr.F32=0.55` reduces F3.2 remaining to 0.19 and raises F3.2 coverage to 0.55. At slot 2, `low_return_alts` scores 1.35 — it addresses F4.6 (return hesitation, still at 0.31) and benefits from the lower competition now that F3.2 is partly addressed. The supporting numbers are at `shapes.F32.*, mods.fit_reassurance.*, mods.low_return_alts.*` in the policy repository.

This answer is **concrete (named contributions with specific values), not narrative (post-hoc rationalisation)**. A SHAP/LIME explanation of a neural ranker is the latter — it is an approximation of how the model behaves locally, computed by an external library. The EDP answer is the policy itself. The reviewer can audit it by reading the same JSON the agent edited last round.

Reviewer-rejection vignette. Edits are not auto-merged. A representative review interaction on the round-2500 edit batch (Fig. 10):

Agent proposes: `easy_returns_promise.on_cov.F46: -0.5 → +0.9`, reason: “NEW synergy: fire AFTER `return_explainer` has built F46 coverage, doubling trust reinforcement for `returner_anxious`.”

Reviewer: “+0.9 is large enough that it could fire even when F4.6 coverage is moderate-low. Could over-trigger for personas that don’t actually need second-layer reassurance. Tighten to +0.5 and add a check that we’re not also penalising it on the F32 axis.”

Agent revises: `easy_returns_promise.on_cov.F46: -0.5 → +0.5`, reason: “Reviewer feedback: smaller positive coverage slope; sufficient for the synergy with `return_explainer` without over-triggering on moderate-F46 sessions.”

This loop — agent proposes, reviewer reads ten lines and writes a one-sentence critique, agent emits a smaller edit — is the closure-under-edits property in action. The representation is closed under both the LLM’s proposal and the human’s revision; neither party needs to translate the other’s language. Neural-weight reviews and SHAP-explanation reviews do not have this property because the language the human reviews (the SHAP plot) is not the language the system updates (the weights).

What the audit cannot do. The trace tells you *what the policy did* and *what the agent changed*. It does not tell you

whether the LLM-authored shape functions encode the right *latent constructs* — the same question one asks of any policy. The taxonomy of 7 problem codes is the LLM’s hypothesis about how customers struggle; if “decision paralysis” is structurally different from “comparison friction” in a way the agent missed, the policy is wrong in a way the audit will not reveal. The closure property is closure on the *operationalisation*, not on the *taxonomy*. We address the taxonomy-discovery question in §3.3 and as an open direction in §8.

0.6.3. When each approach wins

- **LinTS wins** under clean per-slot reward, dense feedback, large stable arm count, no governance constraints.
- **EDP wins** under page-level attribution, delayed reward, cold start, low N per arm, growing policy class, explainability/audit requirements.

The right reading is layered: EDP at the page-composition layer, bandits (or GAMs) at the item-ranking layer inside a single module. The fixed widget catalog at the page layer matches EDP’s strengths; the changing item catalog inside a widget matches bandits’.

Cost. EDP-agent and Bayesian-EDP each consume ~9 subagent calls per 10K-session run (3 checkpoints $\times$ 1 call per round, plus zero per session). At current API rates that is on the order of \$1–3 per run for a frontier model. Per-decision LLM cost is amortised over 2,500 sessions between checkpoints and is essentially zero against the bandit’s per-decision compute cost. The LLM costs are *training-time* costs, not *serving-time*: at serving time EDP is a pure-Python module-scoring function. This is the inverse of the cost profile reviewers might expect from “LLM-in-the-loop” methods.

A fuller picture (numbers from our setup where measured, otherwise approximate for the comparable production stack):

The training-time row is where production budgets are most often surprised. A “neural ranker” deployment that retrains weekly on logged engagement consumes orders of magnitude more compute than the entire LLM-edit loop here, and the resulting weights are not editable by anyone, not even the team that trained them. The serving-time row decides whether the policy can sit inline in a latency-budgeted page render; the dependency row decides whether it survives a model-server outage. EDP wins both rows by construction. The “sessions to cold-start viability” row is the deployment property §5.7 quantifies directly.

0.7. Limitations

1. **Ground-truth circularity is a real risk; we ran the adversarial-perturbation check but not the cross-LLM one.** `TRUE_NEEDS` and `TRUE_PROVISIONS` are LLM-authored, the editor at each checkpoint is a Claude subagent, and the diagnostic report it consumes is computed from rewards generated by those LLM-authored maps. The intentional misalignment be-

	EDP (Bayesian EDP)	CombLinUCB	Neural bandits (MLP)	Neural bandits (GPT)
Deployment complexity	~100k lines of code	~10k lines of code	~10k lines of code	~10k lines of code
Per-decision serving cost	~1 sec CPU (GPT-4o)	~1 sec CPU	~1 sec CPU (GPT-4o)	~1 sec CPU
Estimated agent or testing	~10k lines of code	~10k lines of code	~10k lines of code	~10k lines of code
Scalability to real-world validity	~10k lines of code	~10k lines of code	~10k lines of code	~10k lines of code
Model effectiveness	~10k lines of code	~10k lines of code	~10k lines of code	~10k lines of code
Estimate by a real ML reviewer	~10k lines of code	~10k lines of code	~10k lines of code	~10k lines of code
Overall model quality	~10k lines of code	~10k lines of code	~10k lines of code	~10k lines of code

tween EDP’s internal F-code taxonomy and the 7-need ground truth (§2.2) doesn’t neutralise shared semantic substrate inside the LLM family. Appendix C reports the adversarial perturbation check ($\pm 15\%$ on every provision value, $K = 5$ seeds): the EDP-vs-bandit gap survives, and the relative ranking holds. This rules out the artefactual-numerics worst case, but does **not** rule out a systematic LLM bias affecting both the ground truth and the editor in correlated ways. The cross-LLM check (regenerate `TRUE_PROVISIONS` with GPT-4 / Gemini) remains the right experiment and is the largest validity item open.

- Submodular reward is partly what EDP assumes.** Diminishing returns on (`need` $\times$ `provision`) is the structure adaptive-submodular composition is designed for. GreedyLinTS (§5.4c) inherits the same structural fit. §5.4e adds one non-submodular experiment (substitutes penalty $\alpha = 0.15$ per same-type-pair) and shows the architecture advantage persists: Bayesian-EDP still beats every bandit by 5–15 pp under the non-submodular reward, and the gap *widens* vs the strongest bandit (3 pp $\rightarrow$ 5.5 pp) because page-level attribution makes the cross-slot interaction term even harder for bandits to learn. Other non-submodular regimes (threshold effects, full bilinear interactions, complementarity) are untested.
- Combinatorial bandits tested up to CombLinUCB.** Neural bandits (NeuralUCB, MLP+TS) and counterfactual estimators (IPS, DR) are not in the comparison. The page-attribution finding is structural (slate methods inherit the same credit-assignment problem under page-level reward, §5.3) so we expect a similar pattern, but we have not measured it.
- “Evolvable” is partially aspirational.** §5.4d demonstrated the *capability* to evolve Layer-1 PWL shapes; §5.9 demonstrated the *capability* to add new widgets mid-run. Neither delivered a single-trial empirical win. The demonstrated affordance in this paper is **coefficient updates within a fixed policy class** plus an extensible edit grammar; “evolvable” in the strongest sense (policy-class growth that pays off) remains to be shown with multi-seed evaluation.
- LLM-as-policy (§5.8) is an asymmetric baseline.** The one-shot LLM never sees reward; the agent’s diagnostic report effectively does (via per-persona regret). Framing as “Software 3.0 baseline” understates how much of the gap is the feedback loop vs LLM reasoning. A fairer asymmetry would give the one-shot LLM a summary of observed rewards too.

- Stationary persona distribution between drift tests.** No seasonality, no viral effects. §5.9’s drift test uses a single abrupt shift; richer non-stationarity is unstudied.
- No formal regret bounds.** Treatment is purely empirical.
- Cost is not modeled.** Each agent run consumes K Claude calls per checkpoint $\times$ 3 checkpoints $\times$ R replicates. For a 10K-session simulator at $K=1$, $R=3$ that is ~ 9 subagent invocations per agent variant — a few dollars in API cost, an order of magnitude lower than re-training a neural bandit per checkpoint. A real production budget analysis (per-page or per-decision LLM cost vs the bandit’s pure compute cost) is in §8 future work.
- EDP-agent depends on a capable LLM.** Both the live agent and the OPRO ablation use Claude subagents. Smaller / open-weight LLMs would likely degrade the report-based agent more than OPRO (since the report requires structured reasoning while OPRO is closer to gradient-free pattern matching).
- At $K=3\text{--}4$ reps, several comparisons are underpowered.** The OPRO ablation on parametric (§5.4) is only $\sim 1\sigma$ separated. The Bayesian-EDP margin over EDP-agent on parametric (§5.4b) is small. Per-(persona, category) cell numbers in §5.5 have substantial SE that we don’t quote. Larger K is in §8.

0.8. Future Work

We ran four of the items previously listed here (slate-LinTS in §5.3, drift in §5.9, structural exploration in §5.9, multi-agent ensembles in §5.10) and have a new principled fusion (Bayesian-EDP in §5.4b). The remaining open directions:

Neural bandits. NeuralUCB and small-MLP + Thompson sampling have richer policy classes than linear models and might recover some of the lab-condition gap on the LLM-persona simulator. Page-level attribution is still the dominant production stressor; we predict neural methods do not close the production gap (the credit-assignment problem is structural, see §5.2 and §5.3), but a direct experiment would settle it.

Counterfactual estimators. IPS and Doubly Robust estimators can in principle recover partial per-slot credit if a propensity model is available. These methods need their own exploration policy and a logging policy; integrating them is a non-trivial extension and a separate study.

Layer-1-aware diagnostics + interleaved checkpoints. §5.4d added the Layer-1 PWL-edit capability and observed no single-trial improvement. The most promising follow-ups: (a)

interleave Layer-1-only and Layer-2-only checkpoints rather than mixing both in each batch, so we can attribute marginal value; (b) add Layer-1-specific diagnostics to the checkpoint report (e.g., “X% of high-regret sessions had problem F32 detected at <0.3 despite N1_fit > 0.7”), which would give the agent a sharper signal for when shape edits are warranted. Per-(persona, category) conditional shape functions are a structural step beyond and a separate future item.

Larger ensembles + better validation slicing. §5.10’s 3-draw ensemble is statistically indistinguishable from a single draw under the canonical multi-seed numbers, and per-round diagnostics suggest a structural fragility (validation-slice persona mixture differs from the next live segment) that this K=3 study cannot resolve. Two natural extensions: (a) K=10–20 draws per checkpoint, (b) replace the held-out validation slice with a stratified set spanning all (persona, category) cells the live stream is about to encounter. Both add cost; neither is mechanically difficult.

Live engagement validation. Validate the synthetic-ground-truth ranking against logged engagement data from a deployed system. This is the limitation reviewers will press hardest on; the right way to address it is a logged-eval study, not a richer simulator.

Drift recovery via agent-triggered checkpoints. §5.9’s drift test used a fixed checkpoint schedule. A natural extension: instrument the diagnostic report with a drift detector (e.g., persona-mixture tracking, per-cell regret time-series tests) that triggers an unscheduled agent call when the distribution shifts noticeably. Combined with §5.9’s structural-exploration mechanism, this would close the loop on production non-stationarity.

0.9. Conclusion

Per-slot LinTS is the wrong abstraction for slate-with-submodular-reward problems under page-level reward; the framing was always misapplied. The EDP architecture (adaptive submodular over a shared GAM parameterisation) is the right baseline, and most of the production-stack advantage we measure comes from that architecture — not from the LLM. GreedyLinTS, the EDP architecture with no LLM prior and pure SGD, already closes 6 pp of the lab-vs-prod gap that per-slot LinTS suffers (12.8 % parametric, 13.3 % LLM in production). The LLM agent and continuous SGD then trade off depending on the simulator: **on parametric, EDP-agent (LLM checkpoint edits, no SGD) is the best method we measured (6.5 ± 0.2 %) because the agent’s edits are well-targeted; on LLM-persona, Bayesian-EDP (LLM-anchored prior + continuous regularised SGD) is the best (11.0 ± 0.2 %) because the wider persona / category space leaves enough numerical-calibration headroom for SGD to help. Neither method is uniformly best — they are simulator-conditional. The OPRO ablation isolates the *structured diagnostic report* — not the LLM’s general intelligence — as the carrier of the LLM-side gain, cleanly so on the LLM simulator ($\approx 2.2 \sigma$ separation at K=3-4) and**

weakly so on parametric ($\approx 1 \sigma$ separation, underpowered). The decomposable claim that survives: **under page-level reward, use the adaptive-submodular policy class; if you have an LLM with knowledge of the problem domain, use it as a prior + checkpoint editor; if your simulator is rich enough that the LLM’s discrete edits don’t saturate, add continuous regularised SGD on top.**

0.10. Appendix A: Open-box decision trace

This appendix walks through Fig. 9 in prose. The session is a `size_anxious_new` persona sampled from the parametric simulator at seed 42, browsing the `dresses` category. The whole trace lives in $14 + 7 + 22 = 43$ named numbers; everything below was generated by reading the policy’s internal state, not by any post-hoc interpretability library.

Step 1 — raw signals (14 numbers). The persona’s distribution puts most of its mass on three signals: `size_chart` (0.62), `size_conf` (0.18, so `size_conf_inv` = 0.82), and `return_hist` (0.41). Tab-switch is moderate (0.31), zoom and price-dwell are low. The product-side signal `price_norm` = 0.55 (mid-priced item).

Step 2 — Layer-1 problem fingerprint (7 numbers). The PWL curves in Fig. 8 are evaluated at the signal values above, weighted, and clipped to [0, 1]. The fingerprint comes out concentrated on F3.2:

Problem	Score
F3.2 size anxiety	
F4.6 return hesitation	
F4.1 comparison	
F3.3 / F4.3 / F4.5 / F5.1	

The fingerprint is a 7-d vector with one dominant axis. Layer-2 will see this and compose accordingly.

Step 3 — Layer-2 slot-1 scoring (22 numbers per slot). For slot 1, `remaining` = the fingerprint above, `coverage` = 0. The score decomposes as `base` + $\sum \text{on_rem}[p] \cdot \text{remaining}[p]$ + $\sum \text{on_cov}[p] \cdot \text{coverage}[p]$ − `slot_decay`·0. The top five candidates (Fig. 9, bottom panel):

`fit_reassurance` wins because of the F3.2 `on_rem` slope, not because of a high base. After it is placed, F3.2 `remaining` drops by 0.55 (its `addr` value) to 0.19, F3.2 `coverage` rises by 0.55 to 0.55, and slot 2 scoring reruns with the updated state. `size_guide`, with `on_cov.F32` = −1.4 in the LLM prior, gets *penalised* in slot 2 by $1.4 \times 0.55 = 0.77$ — the explicit anti-synergy stops the page from wasting a second slot on fit guidance.

(The round-2500 agent edit batch in Fig. 10 mutates this exact anti-synergy on `size_guide.on_cov.F32` from −1.4 to +0.6, converting it from a substitute to a chained complement for the size-anxious persona. The reason field

Shape	Size	On remaining	Weight	Category	Value
fit_reassurance	0.45	0.10	0.10	fit_reassurance	0.10
on_rem	0.45	0.10	0.10	on_rem	0.10
size_guide	0.45	0.10	0.10	size_guide	0.10
weight	0.45	0.10	0.10	weight	0.10
val	0.45	0.10	0.10	val	0.10
total	0.45	0.10	0.10	total	0.10

is preserved verbatim. Whether the edit pays off is empirical (§5.4b–5.4c): it does on parametric, less clearly so on LLM.)

What an interpretability reviewer can verify in 5 minutes. That the size-anxious persona’s page leads with `fit_reass` because of the F3.2 on-remaining slope; that the second slot’s identity is sensitive to the `size_guide` / `fit_reassurance` synergy parameter; that swapping the persona changes the fingerprint, the slot-1 winner, and the slot-2 anti-synergy — all by editing values the reviewer can see and write. No SHAP, no LIME, no post-hoc surrogate. The decision trace is the policy.

What an agent can do at the same place. The agent reads exactly the table above (plus regret aggregates) and proposes scalar edits to any of the 43 numbers, each with a free-text reason. Fig. 10 shows what one such batch looks like in practice. The unit of work is a number-with-a-reason, not a weight update.

This is the open-box property the paradigm hinges on. Every other claim — cold-start advantage, audit-trail compatibility, LLM-as-prior fusion — derives from being able to do this on every session.

0.11. Appendix B: Deployment walkthrough

This appendix sketches how the primitive would embed in a production page-render path, so the deployment claims of §1 and §6.3 have a concrete artifact to discuss.

Step 1 — the policy ships as a JSON file. The 245-parameter learnable policy plus the metadata needed for serving (problem labels, widget descriptors, edit-grammar version) is approximately 9 KB of JSON. It is checked into the same repository as the application code. There is no separate model registry, no model server, no feature store.

```
// policy_v_017.json (9.2 KB) { "version": "017",
"approved_by": "alice@team.example", "from_round": 7500, "shapes": { "F32": { "size_chart": {"bps": [0, 0.2, 0.5, 0.8, 1], "vals": [0, 0.1, 0.45, 0.75, 0.95], "weight": 0.45 }, ... }, ... }, "modules": { "fit_reassurance": { "base": 0.10, "addr": { "F32": 0.55, ... }, "on_rem": { "F32": 2.2, ... }, "on_cov": { "F32": -1.2, ... }, ... } }
```

Step 2 — the serving function is ~80 lines of Python. `score_problems`, `score_module`, and `compose` from `edp/policies/edp.py` are the entire decision logic. Given a session feature vector and a category, they return a 6-widget page. No external calls. The same code path runs in CI tests.

```
# Production serving path (sketch) from edp.policy
import compose import json
POLICY = json.load(open("policy_v_017.json"))
```

```
def render_page(session_features, category):
    page = compose(session_features, POLICY["shapes"],
POLICY["modules"]) # page is ['fit_reassurance',
'low_return_alts', ...] return [fill_widget(w,
category) for w in page]
```

`fill_widget(name, category)` pulls cached content from the widget retriever output — already a production capability today. The EDP layer adds about 1 ms to the page render. The 9 KB policy fits in any inline cache and is loaded at process start.

Step 3 — the audit log is the policy repository’s git log. Each agent edit batch is committed as a PR. The PR body is the diagnostic report the agent read plus the edit batch with its reasons. A reviewer reads ten lines and either merges or requests changes (the reviewer-rejection vignette in §6.2 is the realistic shape of one round). The serving policy file is updated by merging the PR; there is no separate deployment.

```
$ git log --oneline policy_v_*.json
b3f2alc policy v017: round-7500 edits --- strength
8d7e120 policy v016: round-5000 edits --- cut over
3a9c2f5 policy v015: round-2500 edits --- revive c
0001abc policy v014: initial LLM-authored config f
```

A new team member reads the git log to understand what the system has learned. A regulator subpoenas the same log and reads the diagnostic reports and reasons. A non-ML PM proposes a manual edit by opening a PR. None of these workflows require ML infrastructure.

Step 4 — Bayesian-EDP adds a small async loop. If the deployment uses Bayesian-EDP, a separate process drains the delayed-reward queue and SGD-updates the policy file every hour or so. The SGD code is ~50 lines (`edp/policies/bayesian_edp.py`). The resulting numeric drift relative to the LLM prior is logged and bounded; if the SGD wants to move a parameter by more than (say) 50 % of its prior value, the move is held for a human checkpoint review (current implementation is unbounded; bounded drift is a one-line addition).

Step 5 — checkpoint cadence is operational. Every K sessions (we use $K = 2,500$ in experiments; production probably 50K–500K depending on traffic), the orchestrator builds the diagnostic report and invokes the agent. This is an offline cron job, not a serving-path dependency. If the agent is unavailable, the previous policy keeps serving indefinitely; the system degrades gracefully.

What this isn’t. This is a *walkthrough*, not a production case study. We have not deployed this stack at Zalando-scale traffic; the production-scale follow-up is left as the next concrete experiment. The artifact-level claims (policy ships as JSON, audit log is git log, no model server) are mechanical

consequences of the primitive’s properties; the latency claim (~1 ms) is measured in-process but not in a production rendering pipeline. The cold-start, audit, and serving-cost claims are first-principles. The benchmark claim is empirical, not a deployment result.

For reproduction commands, code structure, and the interactive demos, see the supplementary material (`supplementary.md`).

0.12. Appendix C: Adversarial-perturbation validity check

We cannot run a true cross-LLM check (no GPT-4 / Gemini access in our pipeline). The next-best validity check is **adversarial perturbation of the LLM-authored `TRUE_PROVISIONS`**: each provision value is multiplied by `uniform(1- $\epsilon$ , 1+ $\epsilon$ )` for $\epsilon = 0.15$ and clipped to $[0, 1]$, with the perturbation seeded independently each replicate. The greedy oracle is re-computed under the perturbed provisions; methods are scored as % of perturbed-oracle reward lost.

Method	Canonical (no perturbation)	Adversarial
EDP-static		10.7 %
EDP-agent		6.5 ± 0.2 %
Bayesian-EDP		7.6 ± 0.3 %
CombLinUCB-warm		10.5 ± 0.1 %
LinTS-warm		18.9 ± 0.1 %

Findings. 1. The EDP-vs-bandit gap survives the perturbation. EDP variants stay in the 6.5–9 % band; the strongest bandit (CombLinUCB) is in the 11.4–13.1 % band. The $\sim 2\times$ margin holds at every seed. **2. The bandits degrade slightly more under perturbation than EDP does** (CombLinUCB +1.7 pp, LinTS +2.8 pp; EDP-agent +0.5 pp). The bandit’s per-session SGD is fitting the noisy *perturbed* reward signal; EDP’s policy class is closer to canonical and absorbs the perturbation more cleanly. **3. Bayesian-EDP becomes the best method under perturbation** (6.5 %, vs EDP-agent at 7.0 %). The SGD’s bias-variance trade-off plays differently when the canonical reward is itself jittered: the regularised drift away from the LLM prior is small enough to absorb perturbation noise but flexible enough to track the new optimum.

What this rules out. It rules out the worst case where EDP’s win was an artefact of the exact numerical values the LLM chose for `TRUE_PROVISIONS`. A ± 15 % perturbation is a meaningful chunk of the $[0, 1]$ provision range; if the win were knife-edge, it would not survive.

What this does not rule out. It does not rule out a *systematic* LLM bias — both `TRUE_PROVISIONS` and the editor are Claude-family. A perturbation that randomly noises individual provision values is structurally different from a different LLM family generating systematically different *patterns* of provision (e.g., a different LLM might think `outfit_completion` strongly addresses N6_trust, not just N5_styling). The true cross-LLM check remains the right experiment for future work, and we list it explicitly in §7 and §8.