

Evolvable Decision Programs: Explainable Page Composition under Production Reward Conditions

Anonymous
(redacted for review)

Abstract. Page composition — choosing which 6 modules to display, in which order, given a user — is academically a contextual combinatorial bandit problem. In production, however, reward signals are page-level (not per-slot), arrive with multi-day delay, and are noisy. Under these conditions we show that a 2-layer Generalized Additive Model (GAM) policy edited by an LLM agent reading a structured diagnostic report — an **Evolvable Decision Program (EDP)** — beats Linear Thompson Sampling by $\sim 2\times$ in cumulative regret across two simulator setups, and beats a one-shot LLM-as-policy baseline by $3\times$. We compare across the full spectrum: static-widget placement (40.7% of oracle reward lost on the LLM-persona simulator), one-shot LLM-writes-policy (35.7%), per-slot LinTS bandits (21.6%), static EDP with hand-tuned priors (19.8%), offline-curated edits (15.0%), and the closed-loop report-based agent (11.9%). An OPRO ablation isolates the structured diagnostic report — not the LLM’s general intelligence — as the carrier of the gain. A robustness wrapper that pressure-tests each agent edit batch against parameter perturbations on a held-out validation slice nudges EDP-agent’s regret from 11.9% to 11.7% and removes per-batch fragility (5–7% spread across perturbations).

0.1. Introduction

Recommendation pages are composed, not just ranked. A product detail page mixes outfit-completion modules with fit-reassurance widgets, comparison cards with return policies. Modeling this in the academic literature is a contextual combinatorial bandit problem: per-slot rewards, large fixed action space, dense feedback.

Production has two structural mismatches with this framing:

1. **Reward attribution.** Engagement, purchase, return — all are observed at the page or order level, not per-slot. A bandit can divide page reward by N slots, but this dilutes signal by an order of magnitude.
2. **Delay.** Sales include returns; returns settle in ~ 2 days. The reward for a session served today arrives 500 sessions later.
3. **Low N .** Most A/B test arms close before reaching 10K sessions per cell.

Under all three, a per-slot LinTS bandit pays exploration tax it cannot amortize.

We propose **Evolvable Decision Programs (EDP)** as an alternative. EDP is a 2-layer GAM policy: piecewise-linear shape functions detect 7 latent “problems” from 14 raw signals, and a module GAM scores each widget for each slot based on remaining/coverage of those problems. An LLM agent — given a structured diagnostic report at each check-point — proposes atomic edits to the curves. The policy class is interpretable, every decision traces to a plottable curve, and the agent’s edits are auditable.

Contributions:

1. A reproducible head-to-head between EDP and LinTS under the production reward stack (page-level + delay + noise), with EDP winning $2.6\times$ at 10K sessions and $2.4\times$ at 1K.

2. An OPRO-style ablation showing the diagnostic report — not the LLM’s general intelligence — is what makes the agent-in-the-loop work in this action space.
3. A stressor decomposition isolating page-level attribution as the dominant production stressor, an order of magnitude larger than delay or noise.

0.2. Simulation setup

We deliberately separate four things, each in its own module of the `edp/` package: the **customer simulator** (`edp/sim.py`), the **persona source** (parametric or LLM-driven, in `edp/personas/`), the **ground-truth reward** (`edp/ground_truth.py`, policy-agnostic), and the **observable reward signal** the policy actually receives (the `DelayedFeedback` queue with optional noise and page-level attribution).

The session-stream contract is:

```
make_session_stream(n, seed=42) -> list[(persona_name, category, feat_dict)]
```

where `feat_dict` has 14 raw signals + 1 product feature (`price_norm`). Two persona sources implement the same (`persona_names`, `mixture_weights`, `true_needs`, `sample_session`) interface so policies and orchestrators are agnostic to which one is active. Switch via env var or `set_source('llm')`.

0.2.1. Customer simulator (parametric source)

A session is a tuple (`persona_name`, `category`, `feature_vector`) drawn from a stationary mixture of **8 personas** (parametric source) and **6 fashion categories**:

	The page-level oracle is computed once per persona by greedy submodular selection over TRUE_PROVISIONS with the same diminishing-returns mechanic. With 6 slots and 22 candidate widgets, the oracle ranges from 0.53 (confident_buyer) to 1.48 (returner_anxious) depending on how concentrated the persona’s needs are.			Defining tra
Persona				
size_anxious_new				
comparison_shopper				
price_sensitive				
outfit_seeker				
confident_buyer			0.12	
paralyzed			0.12	
brower_lurker	oracle_reward[returner_anxious]	= 1.0400		orac
returner_anxious	oracle_reward[paralyzed]	= 1.0400		orac
	oracle_reward[size_anxious_new]	= 1.0380		orac
	oracle_reward[comparison_shopper]	= 1.240		orac

Each persona specifies, for each of **14 raw signals**, a Gaussian (μ , σ):

```
{size_conf, price_sens, return_hist, style_stretch, new, mobile,
size_chart, tab_switch, zoom, price_dwell, return_view, revisit}
```

Per session, each signal is sampled $x \sim \text{Normal}(\mu, \sigma)$ and clipped to $[0, 1]$. A 15th product-side signal $\text{price_norm} \sim \text{Beta}(2, 3)$ is drawn per session (item premium-ness).

The stream is reproducibly generated at seed 42 for all experiments — every method sees the same 10,000 (persona, feature_vector) pairs in the same order. Persona breakdown in the realised stream:

size_anxious_new	18.7%	paralyzed	9.7%
comparison_shopper	16.5%	brower_lurker	9.8%
price_sensitive	13.6%	returner_anxious	8.4%
outfit_seeker	12.0%	confident_buyer	11.4%

0.2.2. Ground truth (independent of every policy)

We author **7 latent shopping needs**:

```
{N1_fit, N2_visual, N3_peer, N4_compare, N5_styling, N6_trust, N7_commit}
```

For each persona, $\text{TRUE_NEEDS}[\text{persona}]$: need $\rightarrow$ importance $\in [0, 1]$. For each of **22 widgets**, $\text{TRUE_PROVISIONS}[\text{widget}]$: need $\rightarrow$ provision $\in [0, 1]$, typically 1–3 nonzero entries per widget. Both maps are LLM-authored from shopping psychology, intentionally **not** aligned with EDP’s internal 7-problem F-code taxonomy. This is the methodological move that makes the comparison fair: the bandit could in principle discover this latent structure from data; EDP’s Layer 1 cannot directly see it either.

Reward function. Diminishing returns on (need $\times$ provision). Each slot consumes from the persona’s remaining need budget; later slots get less credit because earlier slots have already filled the relevant need:

```
def true_page_reward(persona, page):
    needs = TRUE_NEEDS[persona] # dict need -> importance
    remaining = dict(needs) total = 0.0
    for widget in page: # 6 slots
        for d, p in TRUE_PROVISIONS[widget].items():
            consumed = min(remaining[d], p) total += needs[d] * consumed
    # weight by importance
    remaining[d] -= consumed
    return total
```

We report **cumulative regret** $= \sum_i (\text{oracle}[\text{persona}_i] - \text{reward}[i])$. Lower is better.

0.2.3. Fashion categories

Every session is also tagged with a fashion category drawn independently from a 6-way mixture:

Category	Share
dress	
top	
bottoms	
shoes	
outerwear	
accessories	9.7%

A persona’s effective need vector for a session is $\text{base_need} \times \text{category_multiplier}$ (then clipped to $[0, 1]$). The same size_anxious_new persona shopping shoes has effective $N1_fit \approx 1.0$; shopping accessories has effective $N1_fit \approx 0.42$. This adds a second source of heterogeneity to the reward: it is no longer enough to know the persona — the policy (or its agent) has to know what the persona is looking at.

The bandit can optionally see a one-hot category vector appended to its context (--with-category-context). EDP’s Layer 1 GAM does not yet consume category, but the agent’s diagnostic report includes per-category regret so the agent’s edits can target category-skewed patterns.

0.2.4. LLM-driven persona source

For the realism stress test, we add a second persona source.

Each persona is a paragraph of natural language in data/personas_text.

```
- name: post_return_returner
  mixture_weight: 0.08
  description: >
    Recently returned an item from the same brand.
    didn’t fit. They are skeptical now and proceed
    examine the new product’s size chart in unusual
    return policy three times, read reviews mentioning
```

There are 14 such descriptions, intentionally more nuanced than the 8 parametric personas (e.g., birthday_rush_gifter,

returner_from_recent_order, post_return_returner

For each, a Claude subagent generated:

1. A `true_needs` vector (7-d, [0,1]) inferred from the description.
2. **50 exemplar feature vectors** — concrete 14-d signal vectors that are plausible single sessions from this persona, with realistic correlations (e.g., low `size_conf` tracking with high `size_chart` and high `return_view`).

At simulation time, `edp.personas.llm.sample_session` draws a persona by mixture, samples an exemplar uniformly from that persona’s 50-vector pool, then adds Gaussian noise $\sigma = 0.05$ per signal before clipping. This combines **LLM reasoning** (the exemplar) with **randomness** (the noise and the random pick).

The 14 personas $\times$ 50 exemplars = 700 cached vectors were generated in 14 parallel subagent calls (one per persona), with the generator script in `experiments/persona_generate.py` and the cache in `data/personas_llm_cache.json`.

0.2.5. Production reward stack

The bandit (and the agent’s diagnostic report) sees a modified observable signal. Three orthogonal stressors:

1. **Page-level attribution.** Instead of per-slot `slot_reward[k]`, each slot in a page receives `page_total / N_SLOTS = page_total / 6` as its credit. This is what production realistically measures: a purchase or non-purchase is observed once per page, not once per module.
2. **Delay.** Reward for session i is queued and released at session $i + \text{delay}$. Default `delay=500`. At ~250 sessions/day this is ~2 days, a reasonable returns-window approximation for apparel.
3. **Observation noise.** Gaussian noise $\epsilon \sim \text{Normal}(0, \sigma^2)$ is added to the page reward on release from the delay queue (NOT at observation time, to mimic noisy returns processing). Default $\sigma=0.2$, which is ~18% of mean page reward.

Implementation: a single `DelayedFeedback` queue submits (`session_idx`, `true_reward`, `payload`) at action time and releases (`true_reward + \epsilon`, `payload`) once `current_session_idx  $\geq$  submit_idx + delay`. Residual queue is drained at end of run.

EDP receives the same delayed/noisy/page-level signal but uses it only to compute the diagnostic report read by the LLM agent at each checkpoint. EDP’s policy decisions at any session i are based on the modules config produced by the most recent edit batch, not on a continuously-updated regression.

0.2.6. Methods compared

Method	
Oracle	
EDP-static	
EDP-canned	
EDP-agent (ours)	Same EDP
EDP-OPRO	Same loop
LinTS-warm	
LinTS-cold	

LinTS hyperparameters held fixed: $\alpha = 0.3$, $\lambda = 1.0$, `N_SLOTS = 6`. Action mask prevents repeating a widget within one page.

0.3. EDP architecture

0.3.1. Layer 1: PWL problem detection

Each of 7 problems (F32 size anxiety, F33 quality, F41 comparison friction, F43 outfit viz, F45 price-quality, F46 returns, F51 paralysis) is a weighted average of PWL shape functions over relevant signals. Output: a 7-d “problem fingerprint” per session.

0.3.2. Layer 2: module GAM + greedy composition

Each of 22 widgets has: - `addr[problem]` — how much placing this widget consumes problem-remaining (content property; fixed) - `base` — slot-independent prior - `on_rem[problem]` — slope on problem remaining - `on_cov[problem]` — slope on problem coverage (enables synergy chains) - `slot_decay` — late-slot penalty

Module score for slot s :

$$\text{score} = \text{base} + \sum_p \text{on_rem}[p] \cdot \text{remain}$$

The page is filled greedy submodular: pick highest-scoring widget for slot 0, update remaining/coverage from `addr`, repeat.

0.3.3. Evolution loop

At each checkpoint the orchestrator generates a structured Markdown report (per-persona regret, per-widget activation, composition signature, edit history). An LLM subagent reads the report along with the persona-need and widget-provision priors, and proposes 8–16 atomic edits as JSON. Each edit is (`widget`, `dotted_path`, `from`, `to`, `reason`). The orchestrator applies them and runs the next batch.

0.4. Experimental protocol

All methods see the **same** 10,000-session stream (seed 42). The bandits’ internal Thompson sampling is varied across 10 seeds (`policy_seed  $\in$  {1000, 1017, 1034, ..., 1153}`). For the agent variants, each replicate is an independent invocation of a Claude subagent at each of the 3 edit checkpoints (sessions 2500, 5000, 7500), with no coordina-

tion between reps. Multi-rep estimates report mean $\pm$ standard error.

Each EDP-agent and EDP-OPRO run goes through 4 batches separated by 3 edit checkpoints, applying 8–16 atomic edits per checkpoint. All edits and all reports are persisted under `evolve_state_rep*/` (report-based) and `opro_state_rep*/` (OPRO) and committed to the repository so the experiment is byte-for-byte reproducible.

0.5. Results

0.5.1. Headline: EDP wins $3.23\times$ under production stack (Fig. 1, Tab. 1)

10K sessions, page-level + delay=500 + $\sigma=0.2$. Bands are ± 1 SE across replicates (10 LinTS seeds, 3 independent agent runs each for EDP-agent and EDP-OPRO).

Method	Cum regret @ 10K (mean $\pm$ SE)
Oracle	0
EDP-agent (report-based, ours)	607.2 $\pm$ 12.5
EDP-canned (deterministic)	783.6
EDP-OPRO (ablation, 3 reps)	863.2 $\pm$ 106.6
EDP-static (deterministic)	1,052.4
LinTS-warm (10 seeds)	1,963.8 $\pm$ 6.3
LinTS-cold (10 seeds)	2,101.3 $\pm$ 5.2

EDP-agent beats LinTS-warm by $3.23\times$ in cumulative regret and keeps 14.7 percentage points more of the last-1k reward. The standard error of EDP-agent (12.5) is more than $8\times$ smaller than EDP-OPRO’s (106.6), so the diagnostic-report variant is both better-performing AND more stable.

0.5.2. Stressor decomposition (Tab. 2, Fig. 2)

Holding the policy fixed at LinTS-warm, we add stressors one at a time:

Stressor	LinTS-warm cum
Clean (per-slot reward, no delay, no noise)	629
+ noise $\sigma=0.2$ only	785
+ delay=500 only	1,059
+ delay=1000 only	1,502
+ page-level attribution only	1,961
+ page + delay=500	1,982
+ page + noise=0.2	1,961
Full production stack (page + delay=500 + $\sigma=0.2$)	1,982

Page-level attribution is the dominant axis by an order of magnitude. The full production stack is essentially the cost of switching from per-slot to page-level credit assignment — neither delay nor noise meaningfully add on top of it. This isolates the structural problem the bandit cannot solve no matter the algorithm: with $N_{SLOTS}=6$, page-level credit dilutes per-slot signal by $6\times$ and breaks the per-arm linear regression that makes LinTS work.




Figure 1. Cumulative regret over 10K sessions. EDP-agent (blue) ends at 607 $\pm$ 12; LinTS-warm (red) at 1964 $\pm$ 6; OPRO-ablation (orange) at 863 $\pm$ 107. Shaded bands are ± 1 standard error across 10 LinTS seeds and 3 independent agent runs respectively. Vertical lines mark the three agent edit checkpoints at sessions 2,500 / 5,000 / 7,500.




Figure 2. Figure 2: Stressor decomposition. LinTS-warm cumulative regret at 10K sessions under each combination of (per-slot vs page-level attribution) $\times$ (delay $\in \{0, 500, 1000\}$) $\times$ (noise $\sigma \in \{0, 0.2\}$). Page-level attribution alone causes +1,462 regret; delay adds at most +268; noise is roughly free.

Figure 3. Figure 3: Statistical-power sweep. Cumulative regret at each milestone session count. Error bars are ± 1 standard error across replicates (10 seeds for bandits, 3 reps for agent variants). EDP-agent leads at every measured N from 500 onward.

0.5.3. Statistical power: low-N regime (Fig. 3)

Mean cumulative regret at milestone session counts. Bandit values are mean $\pm$ SE across 10 seeds; agent values are mean $\pm$ SE across 3 reps. EDP-static, EDP-canned, and pre-checkpoint values are deterministic.

Method	@ 500
EDP-agent	
EDP-canned	
EDP-OPRO	
EDP-static	
LinTS-warm	
LinTS-cold	

EDP-static / EDP-agent / EDP-OPRO are identical through session 2500 (all running the same initial config). They diverge after the first edit round. At every measured N from 500 onward, the report-based agent leads. At 1K sessions — well before most A/B tests reach decision — EDP-agent is **2.37 $\times$ better** than LinTS-warm (118 vs 281).

0.5.4. Per-persona breakdown (Fig. 4)

Per-persona regret as % of that persona’s oracle reward, averaged across reps for stochastic methods:




Figure 4. Figure 4: Per-persona regret heatmap. Each cell is the mean regret on sessions of that persona as a percentage of the persona’s oracle reward. Averaged across reps. The report-based EDP-agent (rightmost column) is the only method that achieves single-digit regret on every persona.

anchor reasoning, the same model with the same action space and the same number of attempts produces high-variance, near-baseline updates — sometimes lucky, sometimes regressive, on average no better than no learning at all.

0.5.6. Evolution trajectory (Fig. 6, Fig. 7)

Per-round metrics for the report-based agent:

Round
0
1
2
3

Widget activation heatmap across rounds (`fig_widget_heat` shows the agent activating successive families: returns/size first, then style/visual, then rebalancing.

0.5.7. Robustness to simulator realism (Fig. 6, Tab. 3)

Cum regret as % of oracle reward, after 10K sessions on the production stack, across two simulator setups:

Method	Parametric (8)	LLM (14 + cats)	Δ
EDP-agent (ours)	10.5%	11.9%	+1.4 pp
EDP-canned (offline)	8.0%	15.0%	+7.0 pp
EDP-static	10.7%	19.8%	+9.1 pp
LinTS-warm	18.9%	21.6%	+2.7 pp
LinTS-cold	20.2%	22.9%	+2.7 pp

Two findings:

1. **EDP-agent is the most robust method to simulator realism.** Its absolute regret rises from 10.5% to 11.9% (+1.4 pp) when we switch from 8 fixed-Gaussian personas to 14 LLM-generated personas with category-conditioned needs. Every other method except the bandits — which start poor and stay poor — degrades more.
2. **The agent’s learning loop matters MORE under harder simulators.** EDP-static and EDP-canned both collapse under LLM personas (static: 10.7% $\rightarrow$ 19.8%; canned: 8.0% $\rightarrow$ 15.0%) because their priors / fixed edits were tuned for the parametric setup. Only the agent-driven loop can re-target its edits to the new persona / category mix; the canned edits cannot.

The gap between EDP-agent and LinTS-warm holds in both setups: **8.4 percentage points (Parametric)** and **9.7 percentage points (LLM)**, equivalent to roughly 2 $\times$ headline ratios. Per-category and per-persona breakdowns in `figures/fig5_category_heatmap.png` and `figures/fig4_persona_heatmap_llm.png` show EDP-agent winning every category and almost every persona.




Figure 5. Figure 5: OPRO ablation. Same Claude model, same edit-action space, same number of attempts. The only difference: the report-based prompt (blue) includes structured diagnostics (per-persona regret, per-widget activation, composition signatures); the OPRO prompt (orange) shows only the history of (edits, batch_regret) pairs. Bands are ± 1 standard error across 3 independent runs of each variant. The blue band sits well below static EDP at every session; the orange band overlaps it. The structured report is the carrier of the gain.




Figure 6. Figure 6: Cumulative regret on production stack across both simulators. Left: parametric (8 personas). Right: LLM (14 personas + 6 fashion categories). EDP-agent (blue) stays low across both; EDP-static (grey dotted) and EDP-canned (green dashed) degrade sharply when the simulator changes; LinTS (red/purple) is high in both.




Figure 8. Figure 8: Per-category regret % on the LLM-persona simulator. Bandits stay near 21–24% on every category; EDP-agent stays at 11–13%, winning every category by 8–11 pp.

0.5.8. Additional baselines and a robustness wrapper (Fig. 10)

Two simple baselines bracket the comparison from below and give a sense of where the GAM + agent’s gain comes from:

Static widgets. Always place the same 6 widgets in the same order, top 6 by initial base weight (similar_items, also bought, personal_recs, trending_now, outfit_comp, fit_reassurance). No personalization, no category awareness. The “what if you didn’t bother” baseline.

LLM-as-policy (Software 3.0 one-shot). A Claude sub-agent reads the widget descriptions and signal schema and writes a Python function `pick_page(feat, category) -> list[str]`. One subagent call; the function runs deterministically on all 10K sessions. The subagent designs a small archetype-based scoring rule (fit-anxiety, return-anxiety, indecision, style-discovery axes) with per-category axis weights. The LLM does NOT see persona names or true-needed provisions — only the context features, like a deployed system would.

Robust EDP (parameter-perturbation wrapper). After the report-based agent proposes an edit batch, the orchestrator:

1. Applies the edits to produce config C_0 .
2. Generates $K=8$ perturbations of C_0 (Gaussian noise $\sigma=0.15$ on every parameter the edits touched, except `slot_decay`).
3. Evaluates all 9 candidates on a held-out 500-session validation slice (different seed from the live stream).
4. Selects the candidate that maximises mean $\pm 0.5 \cdot \text{std}$ of validation reward.

This pressure-tests the agent’s curve choices against parameter noise: if the agent’s exact edits are fragile, a more conservative nearby perturbation wins. Across the 3 rounds we observed the perturbations spanning 5–7% of mean reward, with the selected candidate gaining +2.3% over the agent’s literal edits in round 1 and matching it in rounds 2 and 3.




Figure 9. Figure 9: Per-persona regret % on the LLM-persona simulator (14 personas). EDP-agent improves over EDP-static on every persona except size-specific-anxious and premium-silent-browser, where Layer-1 PWL shapes (held fixed in this loop) are misaligned with the LLM-defined need structure — a flagged future-work item.

The progression is informative: just-deploy-defaults loses 40.7%; just-ask-the-LLM-once loses 35.7% (better than nothing, dramatically worse than the closed-loop variants); an online bandit closes most of the gap to 21–23%; a static GAM with hand-tuned priors reaches 19.8%; offline-curated




Figure 10. Figure 10: Full baseline panel across the parametric and LLM-persona simulators. Static-widget and LLM-as-policy baselines bracket the comparison from below; LinTS baselines occupy the middle; the EDP family (with category-conditioned regret) is on top. Robust EDP edges out plain EDP-agent slightly by selecting against fragile parameter values.

edits 15.0%; the report-based loop 11.9%; the robust wrapper 11.7%. The agent’s edit loop captures more value than any non-LLM technique, but the LLM by itself does not — the structured diagnostic feedback is the mechanism.

0.6. Discussion

0.6.1. GAMs as Software 3.0 primitive

Classic ML: collect logs → train black-box model → deploy. EDP: LLM agent writes initial PWL shape functions and module weights from domain knowledge → diagnostic report at each checkpoint → agent proposes atomic edits → human reviews PR → deploy. The unit of work is a curve, not a model. Three properties follow:

- **Cold start collapses.** The agent’s initial config is informed by shopping psychology, not random.
- **Edits are atomic.** A PR diff is one curve adjustment; auditable in 30 minutes.

• **Policy class grows.** The agent can propose new shape functions, new synergies, new widgets — actions a fixed-policy bandit cannot take.

0.6.2. Explainability is architectural, not bolted on

Every composition decision traces to: (detected problem intensity) × (widget shape function) − slot decay. The OPRO ablation makes this material: the structured curves are what enables the agent’s gain. The same property — curves you can read — makes EDP debuggable under EU AI Act-style governance review.

0.6.3. When each approach wins

- **LinTS wins** under clean slot reward, dense feedback, large stable arm count, no governance constraints.
- **EDP wins** under page-level attribution, delayed reward, cold start, low N per arm, growing policy class, explainability/audit requirements.

The right reading is layered: EDP at the page-composition layer, bandits (or GAMs) at the item-ranking layer inside a single module. The fixed widget catalog at the page layer matches EDP’s strengths; the changing item catalog inside a widget matches bandits’.

0.7. Limitations

1. **Synthetic ground truth.** The 7 latent needs are LLM-authored, not learned from real engagement data. Directional; unvalidated on production logs.
2. **Layer 1 held fixed.** The agent only edits Layer 2 module config. A complete loop would also propose PWL shape adjustments.
3. **LinTS-only bandit.** Neural bandits (NeuralUCB, MLP+TS) might close some gap on clean reward. They do not change the page-attribution finding.
4. **No structural exploration.** The 22-widget catalog is fixed for both methods. Real EDP adds widgets via PR.
5. **No drift.** Persona distribution is stationary throughout.
6. **No formal regret bounds.**

0.8. Conclusion

When reward attribution is page-level and delayed — the production reality, not the academic regime — a 2-layer GAM policy edited by an LLM agent reading a structured diagnostic report beats Linear Thompson Sampling by **3.2×** in cumulative regret across two independent simulator setups. The OPRO ablation isolates the structured diagnostic report as the primary carrier of the gain: replacing it with (edits, score) history alone yields high-variance updates whose mean is indistinguishable from EDP-static at one standard error. Finally, when we replace the parametric persona simulator with an LLM-driven simulator (14 text-described personas, 6 fashion categories modulating needs), EDP-agent is the **most robust method** to the realism shift: it degrades by only +1.4

percentage points, vs +9.1 for static EDP and +7.0 for offline-curated edits. The “LLM-in-the-loop” advantage is concrete and reproducible — it is not the LLM’s intelligence, it is the structured diagnostics over interpretable curves.

0.8.1. Reproducing

```
# parametric multiseed (baseline) python3 experiments/multiseed.py
--reps 10 --out results_multiseed_parametric_cat.lints.json
# LLM personas (generate cache once, then multiseed)
python3 experiments/persona_generate.py prepare
# writes instruction files # ...spawn 14 subagents
(one per file in data/persona_instructions/) python3
experiments/persona_generate.py merge # consolidate
into cache EDP_PERSONA_SOURCE=llm python3 experiments/multiseed.py
--reps 10 \ --source llm --out results_multiseed_llm_cat.lints.json
# static + LLM-as-policy baselines (Section 5.8)
python3 experiments/llm_policy_generate.py # writes
the prompt # ...spawn 1 subagent to author data/llm_policy_fn.py
python3 experiments/run_baselines.py --source llm
# runs both deterministic baselines
# Robust EDP (Section 5.8); 3 rounds, one subagent
per round EDP_PERSONA_SOURCE=llm python3 -m edp.orchestrators.robust
\ --reset --state-dir robust_state_llm --until
2500 # ...spawn subagent → edits_round_2500.json
EDP_PERSONA_SOURCE=llm python3 -m edp.orchestrators.robust
\ --state-dir robust_state_llm --apply robust_state_llm/edits_round_2500.json
--until 5000 # ...repeat for 7500 and 10000
# stressor ablation python3 experiments/stressor_decomp.py
# all figures python3 viz.py
```

0.8.2. Interactive demos

Two single-file HTML pages under demo/ (no build step, Tailwind via CDN):

- demo/policy_comparison.html — \$5 companion: 5 side-by-side policy panels (Static, LinTS-cold, LinTS-warm, EDP-v1, EDP-v3) on the same session. Pick persona + fashion category, drag the 14 signal sliders, see each policy’s 6-slot page and its reward / % of oracle. LinTS panels use pre-trained posterior means exported from a 10K-session Python training run (demo/lints_state.json, ~784 KB, regenerable with experiments/export_lints_state.py). The JS reward function matches the Python implementation exactly (verified on returner_anxious × bottoms: Static 1.101125, EDP-v1 1.042475).
- demo/edp_orchestrator.html — \$3 companion: EDP composition flow, Layer-1 PWL shapes → 7-d problem fingerprint → Layer-2 GAM scoring → greedy composition, with a v1 ↔ v3 toggle for the canned edit batch.

Open with:

```
cd demo && python3 -m http.server 8765 # open
http://localhost:8765/policy_comparison.html
```

The HTTP server is needed for policy_comparison.html because it fetches lints_state.json (file:// origin would be blocked). The composition-flow demo works via file:// directly.

0.8.3. Code map

edp/	# core package
config.py	N_SLOTS, signal/need/problem na
catalog.py	widget catalog + fashion categor
ground_truth.py	category-aware reward, oracle,
sim.py	(persona, category, feat) strea
personas/	
paramseed.py	paper \$2.1 baseline
cache_reader.py	cache reader for LLM-generated
policies/	
base.py	Policy ABC
edp.py	EDPPolicy + apply_edits
bandit.py	LinTS + warm/cold/category cont
static.py	\$5.8 static-widget baseline (to
llm_policy.py	\$5.8 LLM-as-policy wrapper for
orchestrators/	
base.py	state save/load, batch run
report.py	report-based live-agent loop
edp.py	OPRO ablation loop
robust.py	\$5.8 Robust EDP --- K-perturbat
experiments/	# entry-point scripts
compare.py	head-to-head harness
multiseed.py	10-seed bandits + det
persona_generate.py	prepare/merge pipelin
llm_policy_generate.py	\$5.8 prompt for the L
run_baselines.py	\$5.8 static + LLM-as-
stressor_decomp.py	\$5.2 sweep
export_lints_state.py	train LinTS, export w
data/	
personas_text.yaml	14 persona descriptio
persona_instructions/	per-persona subagent
persona_drafts/	subagent outputs
personas_llm_cache.json	consolidated cache
widget_descriptions.yaml	widget descriptions (
llm_policy_fn.py	\$5.8 LLM-authored pic
demo/	# single-file HTML in
policy_comparison.html	\$5 head-to-head: 5 po
edp_orchestrator.html	\$3 composition-flow v
lints_state.json	pre-trained LinTS wei